

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from scipy.stats import skew

from imblearn.over_sampling import SMOTE

from sklearn.preprocessing import TargetEncoder, StandardScaler
from sklearn.model_selection import train_test_split, KFold, cross_val_score, Gr
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier

from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, roc_auc_s
```

```
In [ ]: df=pd.read_csv("ola_driver_scaler.csv")
```

## Define Problem Statement and perform Exploratory Data Analysis

### Problem Statement

Recruiting and retaining drivers is seen by industry watchers as a tough battle for Ola. Churn among drivers is high and it's very easy for drivers to stop working for the service on the fly or jump to Uber depending on the rates.

As the companies get bigger, the high churn could become a bigger problem. To find new drivers, Ola is casting a wide net, including people who don't have cars for jobs. But this acquisition is really costly. Losing drivers frequently impacts the morale of the organization and acquiring new drivers is more expensive than retaining existing ones.

You are working as a data scientist with the Analytics Department of Ola, focused on driver team attrition. You are provided with the monthly information for a segment of drivers for 2019 and 2020 and tasked to predict whether a driver will be leaving the company or not based on their attributes like

Demographics (city, age, gender etc.) Tenure information (joining date, Last Date)  
Historical data regarding the performance of the driver (Quarterly rating, Monthly business acquired, grade, Income)

### Column Profiling:

- MMMM-YY : Reporting Date (Monthly)
- Driver\_ID : Unique id for drivers
- Age : Age of the driver

- Gender : Gender of the driver – Male : 0, Female: 1
- City : City Code of the driver
- Education\_Level : Education level – 0 for 10+ ,1 for 12+ ,2 for graduate
- Income : Monthly average Income of the driver
- Date Of Joining : Joining date for the driver
- LastWorkingDate : Last date of working for the driver
- Joining Designation : Designation of the driver at the time of joining
- Grade : Grade of the driver at the time of reporting
- Total Business Value : The total business value acquired by the driver in a month (negative business indicates cancellation/refund or car EMI adjustments)
- Quarterly Rating : Quarterly rating of the driver: 1,2,3,4,5 (higher is better)

## Observations in Data

In [ ]: `df.head()`

Out [ ]:

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateof
0	0	01/01/19	1	28.0	0.0	C23	2	57387	2
1	1	02/01/19	1	28.0	0.0	C23	2	57387	2
2	2	03/01/19	1	28.0	0.0	C23	2	57387	2
3	3	11/01/20	2	31.0	0.0	C7	2	67016	1
4	4	12/01/20	2	31.0	0.0	C7	2	67016	1

In [ ]: `df.tail()`

Out [ ]:

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	D
19099	19099	08/01/20	2788	30.0	0.0	C27	2	70254	
19100	19100	09/01/20	2788	30.0	0.0	C27	2	70254	
19101	19101	10/01/20	2788	30.0	0.0	C27	2	70254	
19102	19102	11/01/20	2788	30.0	0.0	C27	2	70254	
19103	19103	12/01/20	2788	30.0	0.0	C27	2	70254	

In [ ]: `# dropping Unnamed: 0 col`  
`df.drop(columns="Unnamed: 0",inplace=True)`

In [ ]: `# descriptive summary`  
`df.describe()`

Out[ ]:

	Driver_ID	Age	Gender	Education_Level	Income	Desi
<b>count</b>	19104.000000	19043.000000	19052.000000	19104.000000	19104.000000	19104
<b>mean</b>	1415.591133	34.668435	0.418749	1.021671	65652.025126	1
<b>std</b>	810.705321	6.257912	0.493367	0.800167	30914.515344	0
<b>min</b>	1.000000	21.000000	0.000000	0.000000	10747.000000	1
<b>25%</b>	710.000000	30.000000	0.000000	0.000000	42383.000000	1
<b>50%</b>	1417.000000	34.000000	0.000000	1.000000	60087.000000	1
<b>75%</b>	2137.000000	39.000000	1.000000	2.000000	83969.000000	2
<b>max</b>	2788.000000	58.000000	1.000000	2.000000	188418.000000	5

- drivers have varied ages from 21 (lowest) to 58 (highest) years of age
- there are 59% male drivers and 41% female drivers

In [ ]:

```
# checking datatypes
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MMM-YY                19104 non-null  object
1   Driver_ID             19104 non-null  int64
2   Age                   19043 non-null  float64
3   Gender                19052 non-null  float64
4   City                  19104 non-null  object
5   Education_Level       19104 non-null  int64
6   Income                19104 non-null  int64
7   Dateofjoining         19104 non-null  object
8   LastWorkingDate       1616 non-null   object
9   Joining Designation   19104 non-null  int64
10  Grade                 19104 non-null  int64
11  Total Business Value  19104 non-null  int64
12  Quarterly Rating      19104 non-null  int64
dtypes: float64(2), int64(7), object(4)
memory usage: 1.9+ MB
```

In [ ]:

```
# Lets check data for a given driverID
df[df['Driver_ID'] == 18]
```

Out[ ]:

	MMM- YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	La
<b>63</b>	01/01/19	18	27.0	1.0	C17	1	31631	01/09/19	
<b>64</b>	02/01/19	18	27.0	1.0	C17	1	31631	01/09/19	
<b>65</b>	03/01/19	18	27.0	1.0	C17	1	31631	01/09/19	
<b>66</b>	04/01/19	18	27.0	1.0	C17	1	31631	01/09/19	
<b>67</b>	05/01/19	18	27.0	1.0	C17	1	31631	01/09/19	

In [ ]:

```
val = df[['Driver_ID']].sample()
did = val.iloc[0,0]
df[df['Driver_ID'] == did]
```

Out[ ]:

	MMM- YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	
<b>13192</b>	01/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13193</b>	02/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13194</b>	03/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13195</b>	04/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13196</b>	05/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13197</b>	06/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13198</b>	07/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13199</b>	08/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13200</b>	09/01/19	1959	35.0	0.0	C26	2	40099	19/06/18	
<b>13201</b>	10/01/19	1959	36.0	0.0	C26	2	40099	19/06/18	
<b>13202</b>	11/01/19	1959	36.0	0.0	C26	2	40099	19/06/18	
<b>13203</b>	12/01/19	1959	36.0	0.0	C26	2	40099	19/06/18	
<b>13204</b>	01/01/20	1959	36.0	0.0	C26	2	40099	19/06/18	
<b>13205</b>	02/01/20	1959	36.0	0.0	C26	2	40099	19/06/18	

In [ ]:

```
# converting objects to datetime format
df['Dateofjoining']=pd.to_datetime(df['Dateofjoining'],errors='coerce')
df['LastWorkingDate']=pd.to_datetime(df['LastWorkingDate'],errors='coerce')
```

C:\Users\ashut\AppData\Local\Temp\ipykernel\_32560\402060823.py:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df['Dateofjoining']=pd.to_datetime(df['Dateofjoining'],errors='coerce')
```

C:\Users\ashut\AppData\Local\Temp\ipykernel\_32560\402060823.py:3: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df['LastWorkingDate']=pd.to_datetime(df['LastWorkingDate'],errors='coerce')
```

```
In [ ]: df[df['Driver_ID'] == 18]
```

Out[ ]:

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	La
63	01/01/19	18	27.0	1.0	C17	1	31631	2019-01-09	
64	02/01/19	18	27.0	1.0	C17	1	31631	2019-01-09	
65	03/01/19	18	27.0	1.0	C17	1	31631	2019-01-09	
66	04/01/19	18	27.0	1.0	C17	1	31631	2019-01-09	
67	05/01/19	18	27.0	1.0	C17	1	31631	2019-01-09	

```
In [ ]: # grouping over driverid to have our data behaviour consistent
```

```
def check(x):
    val = x.unique()
    return val[-1]

groups = df.groupby("Driver_ID")
# agg_dict = {
#     "MMM-YY" : "count",
#     "Age" : 'first',
#     "Gender" : "first",
#     "City" : 'first',
#     "Education_Level" : 'first',
#     "Income" : 'first',
#     "Dateofjoining" : 'first',
#     "LastWorkingDate" : 'last',
#     "Joining Designation" : 'first',
#     "Grade" : 'first',
#     "Total Business Value" : "sum" ,
#     "Quarterly Rating" : 'mean'
# }
```

```
dfg = groups.agg(
    total_months=("MMM-YY" , "count"),
    age=('Age', 'first'),
    gender=('Gender', "first"),
    city=("City", 'first'),
    education=("Education_Level" , "first"),
    income=("Income", 'last'),
    DOJ= ("Dateofjoining" , 'first'),
    LWD= ("LastWorkingDate" , check) ,
    joining_designation=("Joining Designation" , 'first'),
```

```

grade=("Grade" , 'first'),
total_business_val=("Total Business Value" , "sum"),
quarterly_rating=("Quarterly Rating" , 'mean')
).reset_index()

```


```
In [ ]: print(f"Total number of unique driverID : {groups.ngroups}")
```

Total number of unique driverID : 2381

```
In [ ]: dfg.head()
```

```
Out[ ]:   Driver_ID  total_months  age  gender  city  education  income  DOJ  LWD  joinin
```

0	1	3	28.0	0.0	C23	2	57387	2018-12-24	2019-03-11
1	2	2	31.0	0.0	C7	2	67016	2020-11-06	NaT
2	4	5	43.0	0.0	C13	2	65603	2019-12-07	2020-04-27
3	5	3	29.0	0.0	C9	0	46368	2019-01-09	2019-03-07
4	6	5	31.0	1.0	C11	1	78728	2020-07-31	NaT



```
In [ ]: # creating the target
dfg['Churn']=dfg['LWD'].fillna(0).apply(lambda x : 1 if x!=0 else 0)
dfg.drop(columns='LWD',inplace=True)
```

```
In [ ]: dfg.head()
```

```
Out[ ]:   Driver_ID  total_months  age  gender  city  education  income  DOJ  joining_desig
```

0	1	3	28.0	0.0	C23	2	57387	2018-12-24	
1	2	2	31.0	0.0	C7	2	67016	2020-11-06	
2	4	5	43.0	0.0	C13	2	65603	2019-12-07	
3	5	3	29.0	0.0	C9	0	46368	2019-01-09	
4	6	5	31.0	1.0	C11	1	78728	2020-07-31	



```
In [ ]: dfg['gender'].value_counts(normalize=True)
```

```
Out[ ]: gender
0.0    0.589668
1.0    0.410332
Name: proportion, dtype: float64
```

```
In [ ]: # target distribution
dfg['Churn'].value_counts(normalize=True)
```

```
Out[ ]: Churn
1    0.678706
0    0.321294
Name: proportion, dtype: float64
```

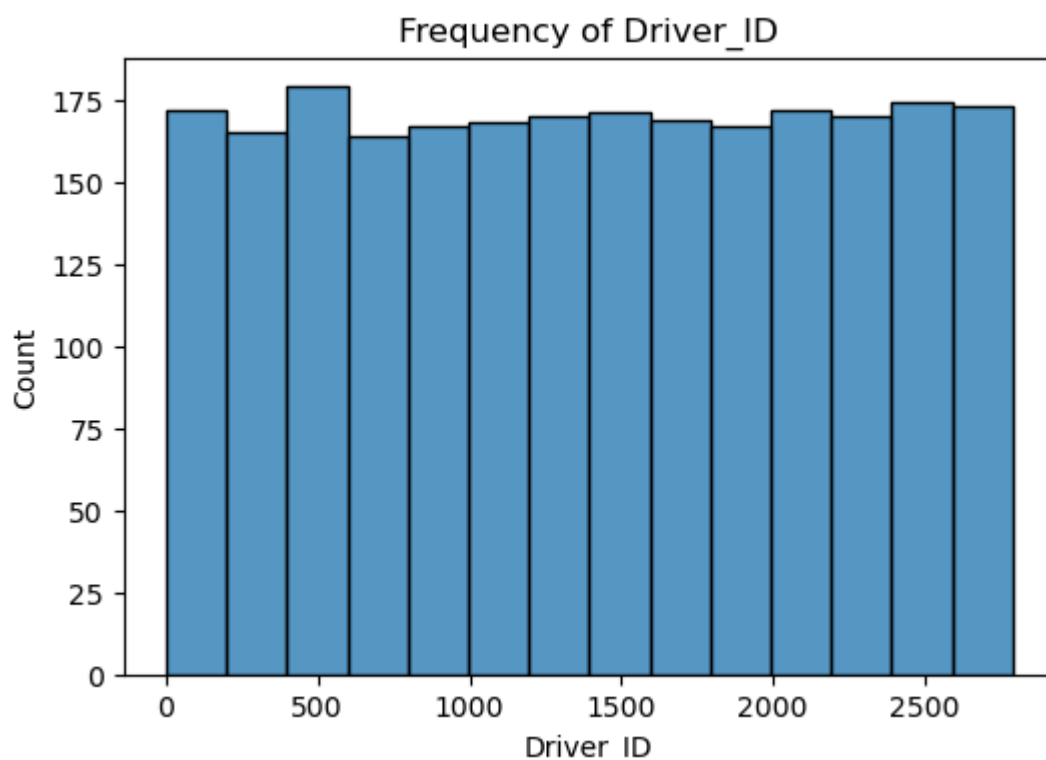
## Univariate Analysis

```
In [ ]: continuous_cols = dfg.select_dtypes(exclude='object').columns
categorical_cols = dfg.select_dtypes(include='object').columns
```

```
In [ ]: # plotting histogram for continuous values
continuous_cols
```

```
Out[ ]: Index(['Driver_ID', 'total_months', 'age', 'gender', 'education', 'income',
              'DOJ', 'joining_designation', 'grade', 'total_business_val',
              'quarterly_rating', 'Churn'],
              dtype='object')
```


```
In [ ]: col = "Driver_ID"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```



data is uniformly distributed

```
In [ ]: col = "total_months"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
```


```
plt.title(f"Frequency of {col}")  
plt.show()
```



data is right skewed, most of the drivers have records for 0-5 months [i.e most of the driver have worked 0-5 months]


```
In [ ]: col = "age"  
plt.figure(figsize=(6,4))  
sns.histplot(data=dfg,x=col)  
plt.title(f"Frequency of {col}")  
plt.show()
```





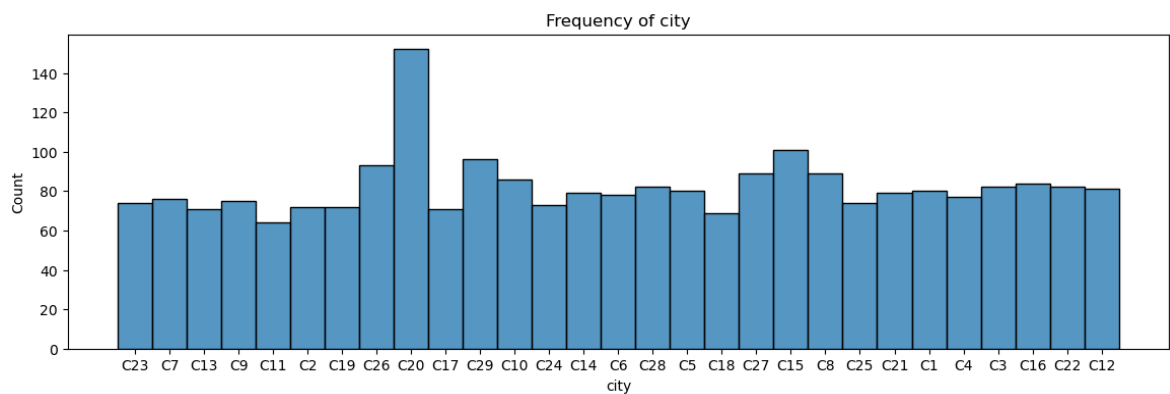
data is almost normally distributed, with most age between 33-35

```
In [ ]: col = "gender"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```




there are more male (0) than female (1) drivers in the data

```
In [ ]: col = "city"
plt.figure(figsize=(14,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```




data is uniformly distributed with city with citycode c20 having the maximum frequency in the data

```
In [ ]: col = "education"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```



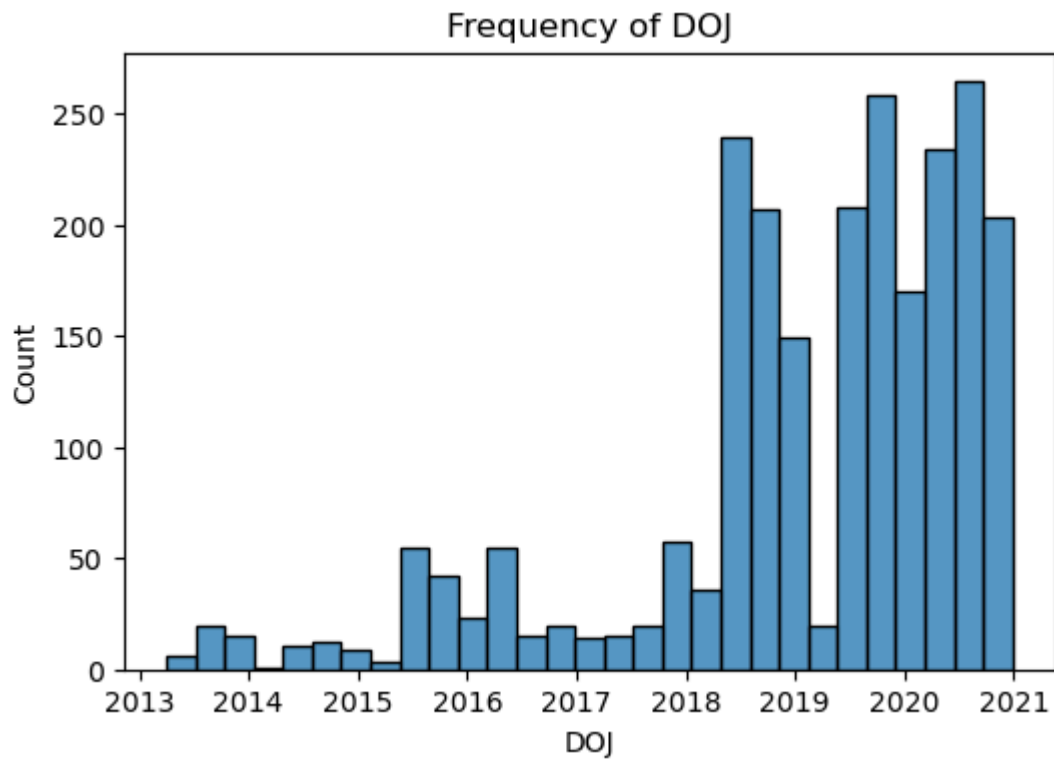
education data is uniformly distributed ( we have diverse drivers with different educational backgrounds)

```
In [ ]: col = "income"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```



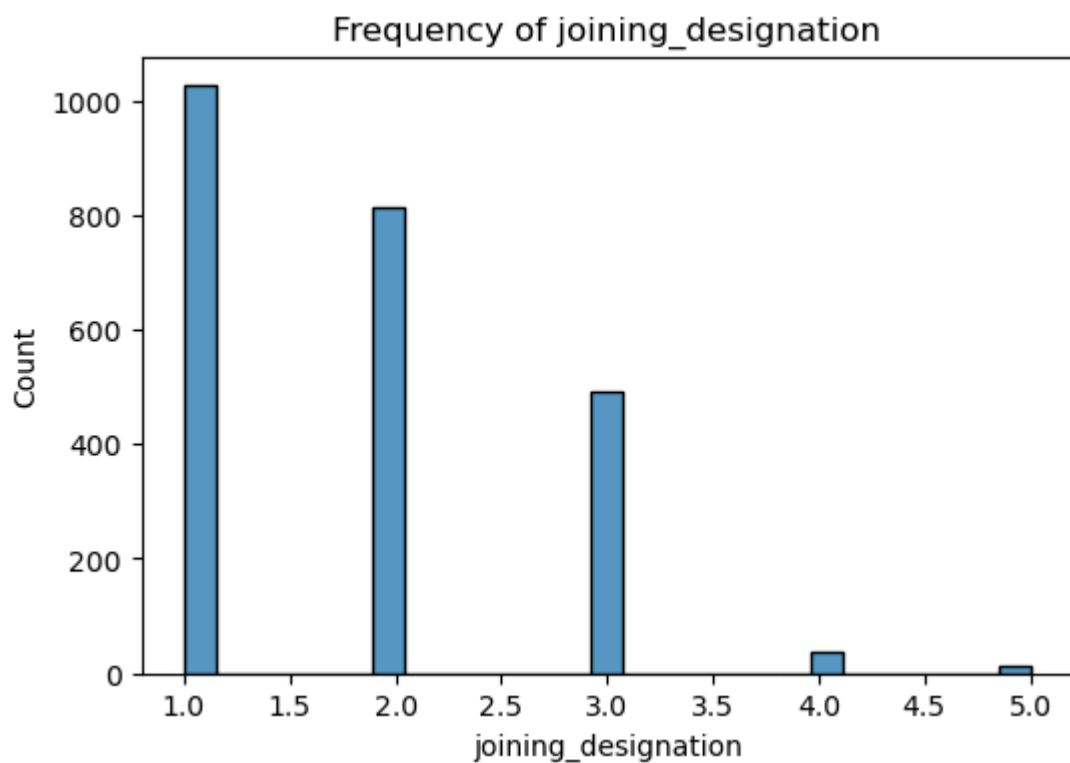
data is having a right skewness with most drivers generating income between 50,000 - 70,000

```
In [ ]: col = "DOJ"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```




DOJ is having left skewness; most of the drivers in the company have joined in recent years [between 2018 - 2021]

```
In [ ]: col = "joining_designation"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```



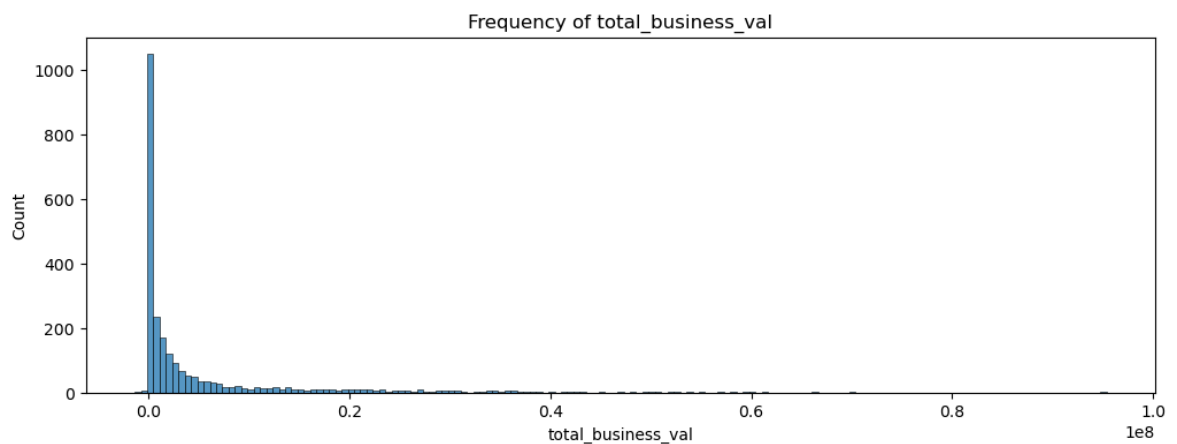
data have a right skewness; most of the drivers when they start working their joining designation is 1.0

```
In [ ]: col = "grade"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```



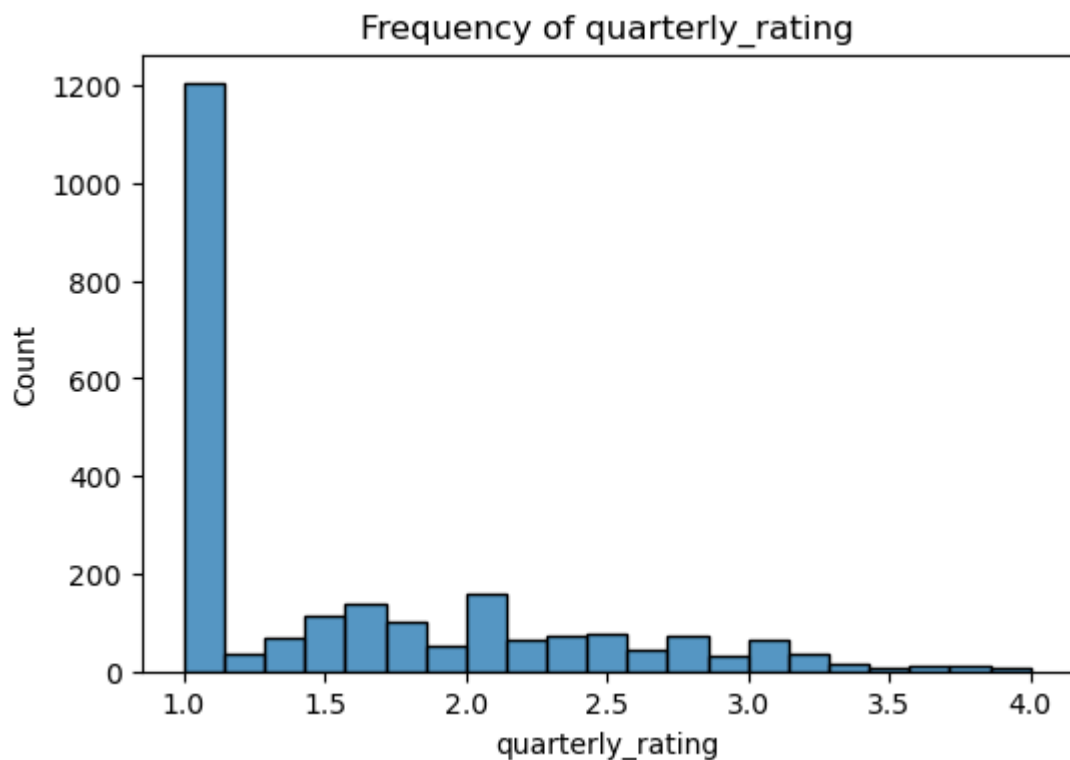
data have right skewness with most drivers having grade of 2.0

```
In [ ]: col = "total_business_val"
plt.figure(figsize=(12,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```




income data is extremely right skewed

```
In [ ]: col = "quarterly_rating"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```



data is having right skewness, with most driving receiving a quarter rating of 1.0

```
In [ ]: col = "Churn"
plt.figure(figsize=(6,4))
sns.histplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```




data shows that most of the drivers have churned as compared to non-churn drivers.

```
In [ ]: # plotting histogram for contiuous values
categorical_cols
```

```
Out[ ]: Index(['city'], dtype='object')
```

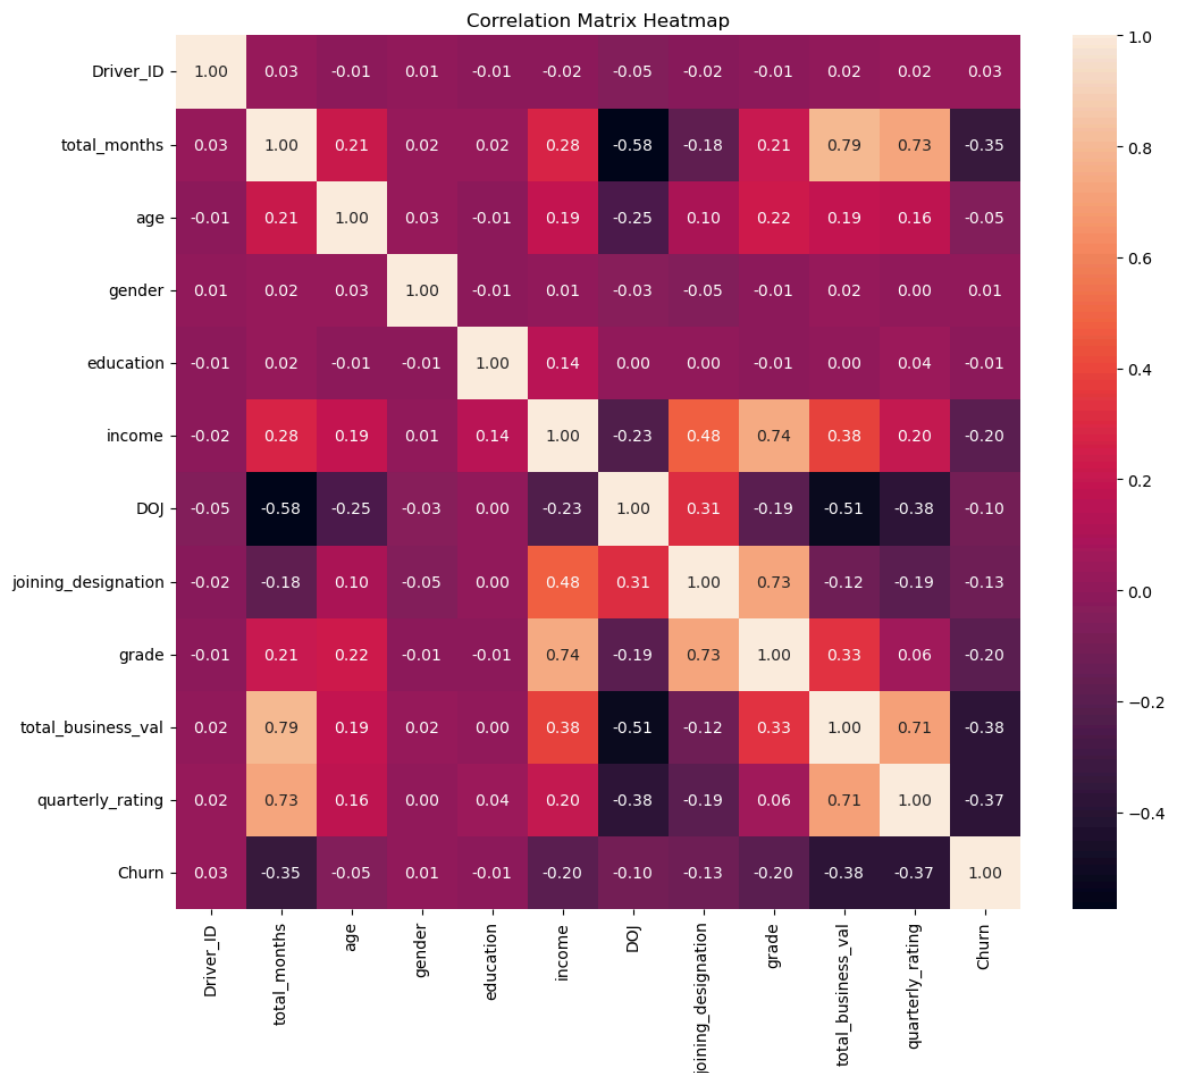
```
In [ ]: col = "city"
plt.figure(figsize=(12,6))
sns.countplot(data=dfg,x=col)
plt.title(f"Frequency of {col}")
plt.show()
```



city with citycode : C20 has the highest requery in the data, this shows that this city could be a hotstop for the ride-sharing company

## Bivariate Analysis

```
In [ ]: # checking correlation of numerical cols
plt.figure(figsize=(12,10))
corr_matrix=dfg[continuous_cols].corr()
sns.heatmap(corr_matrix, annot=True, fmt='.2f')
plt.title('Correlation Matrix Heatmap')
plt.show()
```



```
In [ ]: # Findinf strong correlations
threshold = 0.7

strong_positive_corrs = []
strong_negative_corrs = []

for i in range(len(corr_matrix.columns)):
    for j in range(i):
        if corr_matrix.iloc[i, j] > threshold:
            strong_positive_corrs.append((corr_matrix.columns[i], corr_matrix.co
        elif corr_matrix.iloc[i, j] < -threshold:
            strong_negative_corrs.append((corr_matrix.columns[i], corr_matrix.co
```



```

print(f"Strong Positive Correlations (threshold > {threshold}):")
for col1, col2, corr in strong_positive_corrs:
    print(f"{col1} and {col2}: {corr:.2f}")
print(20*"---")
print(f"Strong Negative Correlations (threshold < -{threshold}):")
for col1, col2, corr in strong_negative_corrs:
    print(f"{col1} and {col2}: {corr:.2f}")

```

Strong Positive Correlations (threshold > 0.7):

grade and income: 0.74  
 grade and joining\_designation: 0.73  
 total\_business\_val and total\_months: 0.79  
 quarterly\_rating and total\_months: 0.73  
 quarterly\_rating and total\_business\_val: 0.71  
 -----

Strong Negative Correlations (threshold < -0.7):

```

In [ ]: # Let change the threshold and check
        threshold = 0.5

        strong_positive_corrs = []
        strong_negative_corrs = []

        for i in range(len(corr_matrix.columns)):
            for j in range(i):
                if corr_matrix.iloc[i, j] > threshold:
                    strong_positive_corrs.append((corr_matrix.columns[i], corr_matrix.co
                elif corr_matrix.iloc[i, j] < -threshold:
                    strong_negative_corrs.append((corr_matrix.columns[i], corr_matrix.co

        print(f"Strong Positive Correlations (threshold > {threshold}):")
        for col1, col2, corr in strong_positive_corrs:
            print(f"{col1} and {col2}: {corr:.2f}")
        print(20*"---")
        print(f"Strong Negative Correlations (threshold < -{threshold}):")
        for col1, col2, corr in strong_negative_corrs:
            print(f"{col1} and {col2}: {corr:.2f}")

```

Strong Positive Correlations (threshold > 0.5):

grade and income: 0.74  
 grade and joining\_designation: 0.73  
 total\_business\_val and total\_months: 0.79  
 quarterly\_rating and total\_months: 0.73  
 quarterly\_rating and total\_business\_val: 0.71  
 -----


Strong Negative Correlations (threshold < -0.5):

DOJ and total\_months: -0.58  
 total\_business\_val and DOJ: -0.51

```

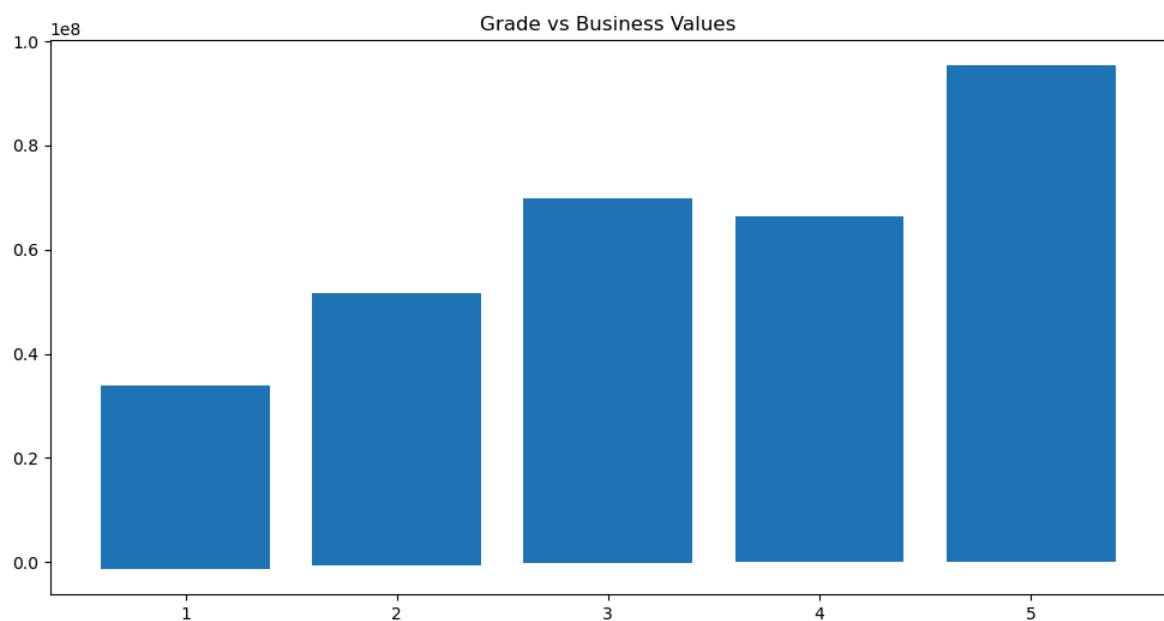
In [ ]: plt.figure(figsize = (12,6))
        plt.bar(dfg['city'], dfg['total_business_val'])
        plt.title("City vs Business Values")
        plt.show()

```




we can see city with citycode : C13 has the highest business value associated with it.

```
In [ ]: plt.figure(figsize = (12,6))
plt.bar(dfg['grade'] , dfg['total_business_val'])
plt.title("Grade vs Business Values")
plt.show()
```



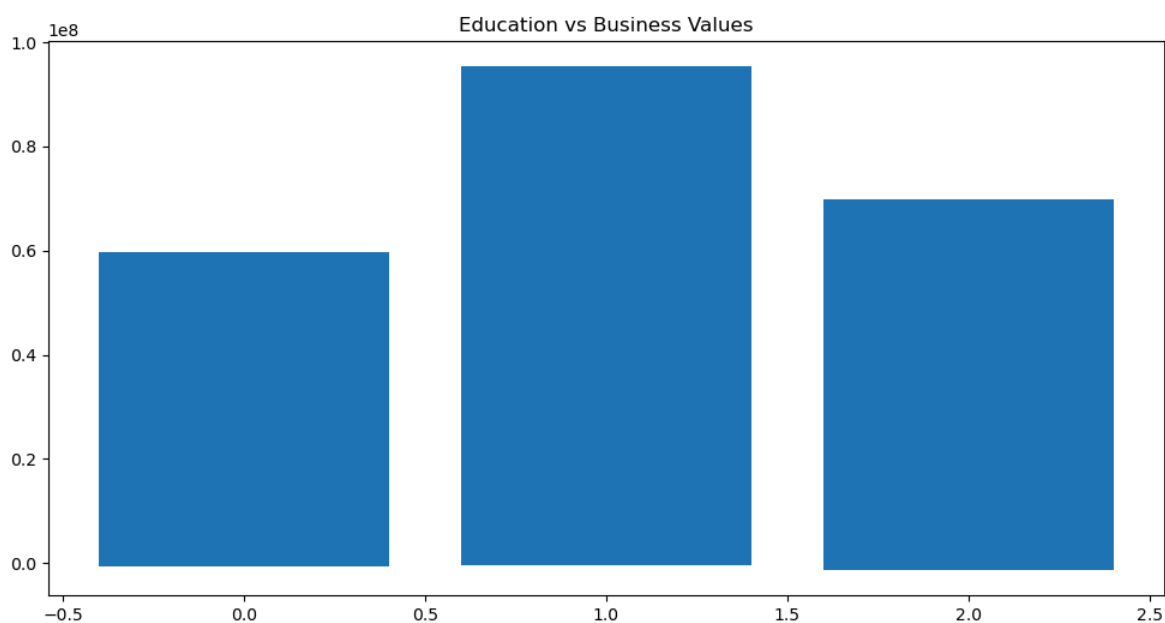
drivers with grade 5 generated highest business values

```
In [ ]: plt.figure(figsize = (12,6))
plt.bar(dfg['quarterly_rating'] , dfg['total_business_val'])
plt.title("Quarterly rating vs Business Values")
plt.show()
```



quarterly rating above 2.5 brings the highest business value


```
In [ ]: plt.figure(figsize = (12,6))
plt.bar(dfg['education'] , dfg['total_business_val'])
plt.title("Education vs Business Values")
plt.show()
```



data shows that drivers with education level : 1 (i.e 12th plus) generates the highest business value followed by education level : 2 (i.e graduate).

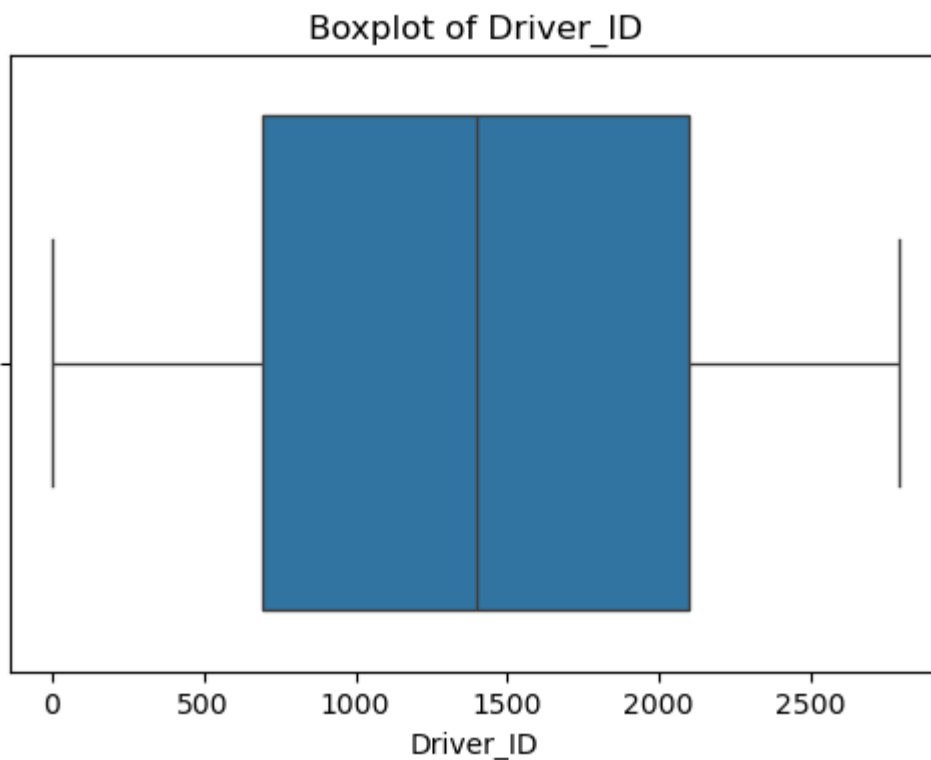
```
In [ ]: # plotting boxplots for categorical col
cat_col='city'
num_col='total_business_val'
# for col in continuous_cols:
plt.figure(figsize = (12,6))
sns.boxplot(data=dfg,x=cat_col, y=num_col)
```

```
plt.title(f"{cat_col} vs {num_col}")
plt.show()
```




city with citycode C13 has the highest outlier in the data


```
In [ ]: # outlier detection and treatment
for col in continuous_cols:
    plt.figure(figsize=(6, 4))
    sns.boxplot(x=dfg[col])
    plt.title(f'Boxplot of {col}')
    plt.show()
```




Boxplot of total\_months




Boxplot of age





Boxplot of gender




Boxplot of education






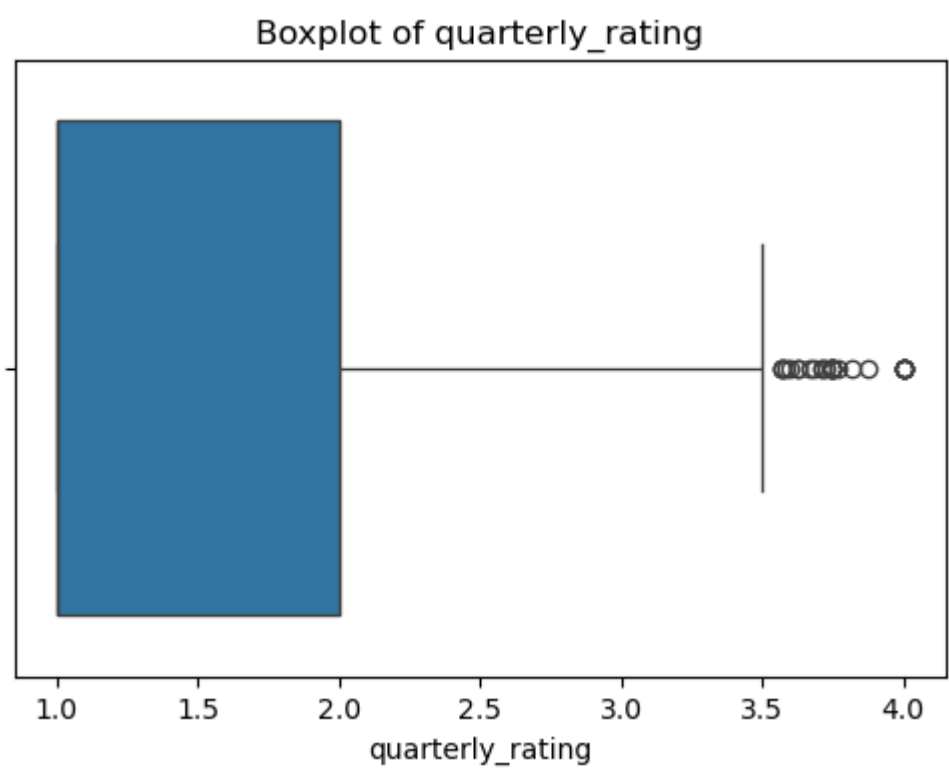
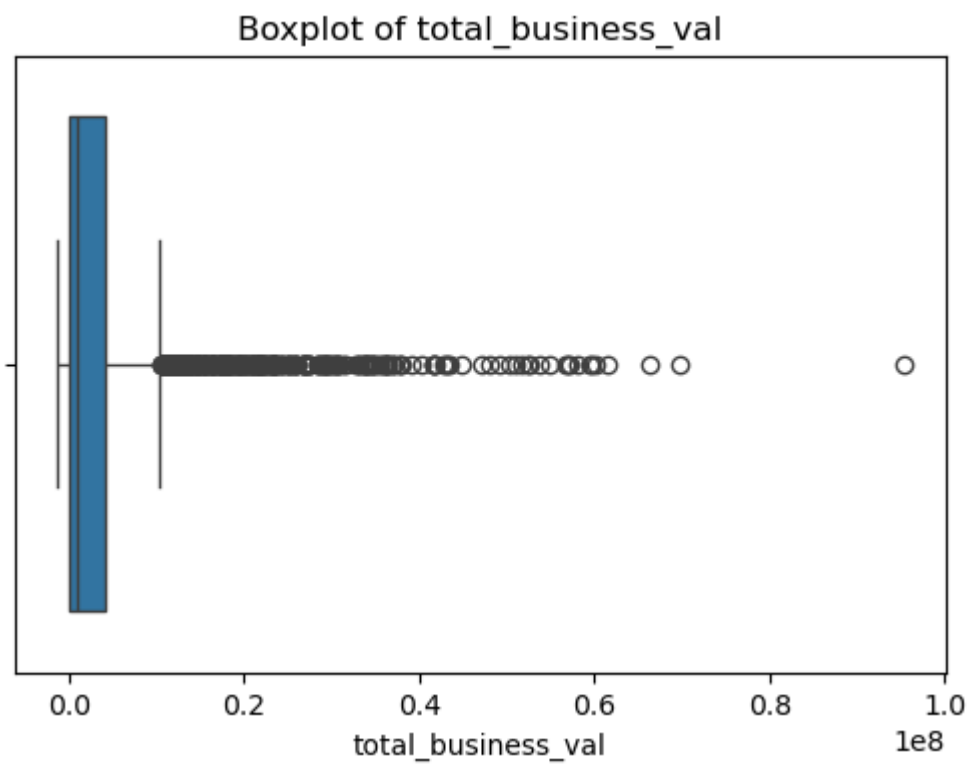
Boxplot of joining\_designation




Boxplot of grade













from visual representation we can see there are columns which have outliers.



```
In [ ]: # treatment
# median imputation : as it preserve the central tendency of the distribution.
# clipping : to limit their impact without distorting the underlying distributio
# Leaving few cols
def outlier_detection(series):
    q1=series.quantile(0.25)
    q3=series.quantile(0.75)
    iqr=q3-q1
    lower_bound = q1 - 1.5*iqr
    upper_bound = q3 + 1.5*iqr
    if series.name in ['age' , 'income'] :
        median_value = series.median()
        output=series.apply(lambda x: median_value if x < lower_bound or x > upp
    else:
        output=series.clip(lower_bound, upper_bound)
    return output
```



```
In [ ]: # treatment for cols
for col in continuous_cols:
    dfg[col] = outlier_detection(dfg[col])
```

```
In [ ]: # checking putlier treatmtent effects
for col in continuous_cols:
    plt.figure(figsize=(6, 4))
    sns.boxplot(x=dfg[col])
    plt.title(f'Boxplot of {col}')
    plt.show()
```













Boxplot of grade



Boxplot of total\_business\_val





```
In [ ]: # checking skewness of the features
def check_skewness(series):
    if series.dtype != 'int64' and series.dtype != 'float64':
        return
    col_skewness = skew(series)
    if col_skewness < -1:
        comment = "Highly left-skewed"
    elif -1 <= col_skewness < -0.5:
        comment = "Moderately left-skewed"
    elif -0.5 <= col_skewness < 0.5:
        comment = "Approximately symmetric"
    elif 0.5 <= col_skewness < 1:
        comment = "Moderately right-skewed"
```



```

else:
    comment = "Highly right-skewed"
    print(f"{series.name}: Skewness = {col_skewness:.2f} ({comment})")

```

```

In [ ]: for col in continuous_cols:
        check_skewness(dfg[col])

```

```

Driver_ID: Skewness = -0.00 (Approximately symmetric)
total_months: Skewness = 1.06 (Highly right-skewed)
age: Skewness = 0.37 (Approximately symmetric)
gender: Skewness = 0.36 (Approximately symmetric)
education: Skewness = -0.01 (Approximately symmetric)
income: Skewness = 0.55 (Moderately right-skewed)
joining_designation: Skewness = 0.44 (Approximately symmetric)
grade: Skewness = 0.52 (Moderately right-skewed)
total_business_val: Skewness = 1.21 (Highly right-skewed)
quarterly_rating: Skewness = 1.03 (Highly right-skewed)
Churn: Skewness = -0.77 (Moderately left-skewed)

```

## Data Preprocessing

```

In [ ]: # checking duplicate data
dup = dfg[dfg.duplicated()]
if dup.empty:
    print("No duplicate records found.")
else:
    print("Duplicate records found:")
    print(dup)

```

No duplicate records found.

```

In [ ]: # missing value detection and treatment
dfg.isna().sum()

```

```

Out[ ]: Driver_ID      0
total_months      0
age              0
gender           0
city             0
education        0
income           0
DOJ              0
joining_designation  0
grade            0
total_business_val  0
quarterly_rating  0
Churn            0
dtype: int64

```

```

In [ ]: # feature engineering
dfg.head()

```

```
Out [ ]:
```

	Driver_ID	total_months	age	gender	city	education	income	DOJ	joining_desig
0	1	3.0	28.0	0.0	C23	2	57387.0	2018-12-24	
1	2	2.0	31.0	0.0	C7	2	67016.0	2020-11-06	
2	4	5.0	43.0	0.0	C13	2	65603.0	2019-12-07	
3	5	3.0	29.0	0.0	C9	0	46368.0	2019-01-09	
4	6	5.0	31.0	1.0	C11	1	78728.0	2020-07-31	

```
In [ ]: # creating date related fields
dfg['year']=dfg['DOJ'].dt.year
dfg['month']=dfg['DOJ'].dt.month
dfg['day']=dfg['DOJ'].dt.day
dfg.drop(columns='DOJ',inplace=True)
```

```
In [ ]: ## creating bins for ages
# print(f"min age : {dfg['age'].min()} \nmax age : {dfg['age'].max()}")

# dfg['age']=pd.cut(dfg['age'] , bins = 5 , labels=['a','b','c','d','e'])
```

```
In [ ]: df.columns
```

```
Out [ ]: Index(['MMM-YY', 'Driver_ID', 'Age', 'Gender', 'City', 'Education_Level',
            'Income', 'Dateofjoining', 'LastWorkingDate', 'Joining Designation',
            'Grade', 'Total Business Value', 'Quarterly Rating'],
            dtype='object')
```

```
In [ ]: # new feature wrt quarterly rating increase
first_quar=df.groupby(['Driver_ID']).agg({'Quarterly Rating' : "first"})
last_quar = df.groupby(['Driver_ID']).agg({'Quarterly Rating' : "last"})

rating_change = (last_quar['Quarterly Rating'] > first_quar['Quarterly Rating'])
rating_change['Quarterly Rating']=rating_change['Quarterly Rating'].astype(int)

dfg['quarterly_rating_increased']=rating_change['Quarterly Rating']
```

```
In [ ]: # new feature wrt income increase
first_income=df.groupby(['Driver_ID']).agg({'Income' : "first"})
last_income = df.groupby(['Driver_ID']).agg({'Income' : "last"})

income_change = (last_income['Income'] > first_income['Income']).reset_index()
income_change['Income']=income_change['Income'].astype(int)

dfg['income_increased']=income_change['Income']
```

```
In [ ]: dfg.head()
```

Out [ ]:

	Driver_ID	total_months	age	gender	city	education	income	joining_designation
0	1	3.0	28.0	0.0	C23	2	57387.0	1.0
1	2	2.0	31.0	0.0	C7	2	67016.0	2.0
2	4	5.0	43.0	0.0	C13	2	65603.0	2.0
3	5	3.0	29.0	0.0	C9	0	46368.0	1.0
4	6	5.0	31.0	1.0	C11	1	78728.0	3.0

In [ ]:

## Model building

In [ ]: `X = dfg.drop(columns='Churn')`  
`y = dfg['Churn']`

In [ ]: `# splitting the data`  
`X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)`

In [ ]: `X_train.head()`

Out [ ]:

	Driver_ID	total_months	age	gender	city	education	income	joining_designation
2242	2627	2.0	34.0	1.0	C7	0	104058.0	
1474	1730	1.0	42.0	0.0	C9	0	51579.0	
2132	2499	3.0	27.0	0.0	C17	1	75458.0	
1873	2200	20.5	34.0	0.0	C15	1	69756.0	
462	539	1.0	36.0	0.0	C24	0	109296.0	

In [ ]: `X_train.shape`


Out [ ]: (1904, 16)

In [ ]: `#encoding`  
`te_cols = ["city"]`  
  
`te = TargetEncoder()`  
`X_train[te_cols]=te.fit_transform(X_train[te_cols],y_train)`  
`X_test[te_cols]=te.transform(X_test[te_cols])`

In [ ]: `X_train.head()`

Out [ ]:


	Driver_ID	total_months	age	gender	city	education	income	joining_desi
<b>2242</b>	2627	2.0	34.0	1.0	0.673188	0	104058.0	
<b>1474</b>	1730	1.0	42.0	0.0	0.673579	0	51579.0	
<b>2132</b>	2499	3.0	27.0	0.0	0.852101	1	75458.0	
<b>1873</b>	2200	20.5	34.0	0.0	0.656125	1	69756.0	
<b>462</b>	539	1.0	36.0	0.0	0.629148	0	109296.0	



In [ ]: `X_test.head()`

Out [ ]:

	Driver_ID	total_months	age	gender	city	education	income	joining_desi
<b>937</b>	1103	5.0	26.0	0.0	0.711650	2	40318.0	
<b>765</b>	899	3.0	29.0	0.0	0.751551	0	28565.0	
<b>34</b>	46	3.0	36.0	1.0	0.753273	2	42171.0	
<b>480</b>	559	5.0	40.0	0.0	0.720025	0	44342.0	
<b>303</b>	358	14.0	39.0	1.0	0.708502	2	117830.0	



```
In [ ]: # scaling
se = StandardScaler()

X_train = se.fit_transform(X_train)
X_test = se.transform(X_test)
```

```
In [ ]: # handling imbalanced data
sm = SMOTE(random_state=42)

X_train_res, y_train_res = sm.fit_resample(X_train , y_train)
# X_test_res = sm.resample(X_test)
```

```
In [ ]: print('Before OverSampling, the shape of train_X: {}'.format(X_train.shape))
print('Before OverSampling, the shape of train_y: {} \n'.format(y_train.shape))

print("Before OverSampling, counts of label '1': {}".format(sum(y_train == 1)))
print("Before OverSampling, counts of label '0': {}".format(sum(y_train == 0)))
```

Before OverSampling, the shape of train\_X: (1904, 16)

Before OverSampling, the shape of train\_y: (1904,)

Before OverSampling, counts of label '1': 1292

Before OverSampling, counts of label '0': 612

```
In [ ]: print('After OverSampling, the shape of train_X: {}'.format(X_train_res.shape))
print('After OverSampling, the shape of train_y: {} \n'.format(y_train_res.shape))

print("After OverSampling, counts of label '1': {}".format(sum(y_train_res == 1)))
print("After OverSampling, counts of label '0': {}".format(sum(y_train_res == 0)))
```

After OverSampling, the shape of train\_X: (2584, 16)

After OverSampling, the shape of train\_y: (2584,)

After OverSampling, counts of label '1': 1292

After OverSampling, counts of label '0': 1292


## Model 0 : DecisionTreeClassifier

```
In [ ]: # model initialization : finding the optimal value for k
model0 = DecisionTreeClassifier(random_state=7)

cv_scores = []
k_values=range(2, 20)
# selecting best value of k
for k in k_values:
    kf=StratifiedKFold(n_splits=k,shuffle=True,random_state=7)
    scores=cross_val_score(model0, X_train, y_train,cv=kf)
    cv_scores.append(np.mean(scores))

# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, cv_scores, marker='o')
plt.xlabel('Number of Folds (k)')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Cross-Validated Accuracy for Different Values of k')
plt.grid()
plt.show()

# Find the best value of k
best_k = k_values[np.argmax(cv_scores)]
print(f"Best value of k: {best_k} with accuracy: {max(cv_scores):.4f}")
```



Best value of k: 10 with accuracy: 0.8850

```
In [ ]: # training base model with best_k
kfold = StratifiedKFold(n_splits = best_k)

class_weights = compute_class_weight(class_weight='balanced', classes=np.unique(
```

```

class_weights_dict = {i: weight for i, weight in enumerate(class_weights)}

dt_clf = DecisionTreeClassifier(random_state=7, class_weight=class_weights_dict)
cv_results_dt = cross_validate(dt_clf, X_train, y_train, cv = kfold, scoring = '

print(f"K-Fold Accuracy Mean: \n Train: {cv_results_dt['train_score'].mean()*100
print(f"K-Fold Accuracy Std: \n Train: {cv_results_dt['train_score'].std()*100:.

```

K-Fold Accuracy Mean:

Train: 100.00

Validation: 86.87

K-Fold Accuracy Std:

Train: 0.00,

Validation: 1.49

```

In [ ]: # initializing model : selecting best estimators
# kfold = KFold(n_splits = 10)
params = {
    "max_depth" : [3, 5, 7, 10, 15],
    "max_leaf_nodes" : [20, 40, 60, 100],
}
grid_dt = GridSearchCV(estimator = DecisionTreeClassifier(random_state=7),
    param_grid= params,
    scoring = 'accuracy',
    cv = kfold,
    n_jobs=-1
)

grid_dt.fit(X_train, y_train)

```

```

Out[ ]: ▸ GridSearchCV ⓘ ?
▸ best_estimator_: DecisionTreeClassifier
    ▸ DecisionTreeClassifier ⓘ

```


```

In [ ]: model_dt = grid_dt.best_estimator_
y_pred= model_dt.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
    display_labels=model_dt.classes_)

disp.plot()
plt.show()

```




```
In [ ]: dt_report = classification_report(y_test, y_pred, output_dict=True)
```

```
In [ ]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.93	0.89	153
1	0.97	0.93	0.95	324
accuracy			0.93	477
macro avg	0.91	0.93	0.92	477
weighted avg	0.93	0.93	0.93	477


```
In [ ]: importances = model_dt.feature_importances_
indices = np.argsort(importances)[::-1]
names = [X.columns[i] for i in indices]

plt.figure(figsize=(15, 7))
plt.title("Feature Importance")
plt.bar(range(X_train.shape[1]), importances[indices])
plt.xticks(range(X_train.shape[1]), names, rotation=45)
plt.show()
```



```
In [ ]: logit_roc_auc=roc_auc_score(y_test,y_pred)
fpr,tpr,thresholds=roc_curve(y_test,model_dt.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr,tpr,label='Decision Tree Classifier (area = %0.2f)' % logit_roc_auc)
plt.plot([0,1],[0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```






### Model 1 : RandomForestClassifier

```
In [ ]: # initialize model : finding best value for k with RF
model1 = RandomForestClassifier(random_state=7)

cv_scores = []
k_values=range(2, 20)
# selecting best value of k
for k in k_values:
    kf=StratifiedKFold(n_splits=k,shuffle=True,random_state=7)
    scores=cross_val_score(model1, X_train, y_train,cv=kf)
    cv_scores.append(np.mean(scores))

# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, cv_scores, marker='o')
plt.xlabel('Number of Folds (k)')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Cross-Validated Accuracy for Different Values of k')
plt.grid()
plt.show()

# Find the best value of k
best_k = k_values[np.argmax(cv_scores)]
print(f"Best value of k: {best_k} with accuracy: {max(cv_scores):.4f}")
```



Best value of k: 8 with accuracy: 0.9144

```
In [ ]: # initialize model : model with defined class_weight

class_weights = compute_class_weight(class_weight='balanced', classes=np.unique(
class_weights_dict = {i: weight for i, weight in enumerate(class_weights)})

kfold = StratifiedKFold(n_splits=best_k)

rf_clf = RandomForestClassifier(random_state=7, class_weight=class_weights_dict)
cv_results_rf = cross_validate(rf_clf, X_train, y_train, cv=kfold, scoring='accu

print(f"K-Fold Accuracy Mean: \n Train: {cv_results_rf['train_score'].mean()*100
print(f"K-Fold Accuracy Std: \n Train: {cv_results_rf['train_score'].std()*100:.
```

K-Fold Accuracy Mean:

Train: 100.00

Validation: 90.97

K-Fold Accuracy Std:


Train: 0.00,

Validation: 1.53

```
In [ ]: # initialize model : selecting best estimators
params = {
    'n_estimators' : [100,200,300,400],
    'max_depth' : [None,3,5,10],
    'criterion' : ['gini', 'entropy'],
    'bootstrap' : [True, False],
    'max_features' : [8,9,10]
}
grid_rf = GridSearchCV(estimator = RandomForestClassifier(random_state=7, class_
    param_grid = params,
    scoring = 'accuracy',
    cv = kfold ,
    n_jobs=-1
)

grid_rf.fit(X_train, y_train)
```

Out[ ]:




```
In [ ]: print("Best params: ", grid_rf.best_params_)
        print("Best score: ", grid_rf.best_score_)
```

```
Best params: {'bootstrap': True, 'criterion': 'entropy', 'max_depth': 10, 'max_f
eatures': 8, 'n_estimators': 100}
Best score: 0.9191176470588236
```

```
In [ ]: model_rf1 = grid_rf.best_estimator_
        y_pred = model_rf1.predict(X_test)

        cm = confusion_matrix(y_test, y_pred)
        disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                                       display_labels=model_rf1.classes_)

        disp.plot()
        plt.show()
```




```
In [ ]: rf_model1=classification_report(y_test, y_pred,output_dict=True)
```

```
In [ ]: print(classification_report(y_test, y_pred))
```


	precision	recall	f1-score	support
0	0.91	0.93	0.92	153
1	0.97	0.96	0.96	324
accuracy			0.95	477
macro avg	0.94	0.94	0.94	477
weighted avg	0.95	0.95	0.95	477

```
In [ ]: importances = model_rf1.feature_importances_
indices = np.argsort(importances)[::-1]
names = [X.columns[i] for i in indices]

plt.figure(figsize=(15, 7))
plt.title("Feature Importance")
plt.bar(range(X_train.shape[1]), importances[indices])
plt.xticks(range(X_train.shape[1]), names, rotation=45)
plt.show()
```



```
In [ ]: logit_roc_auc=roc_auc_score(y_test,y_pred)
fpr, tpr, thresholds=roc_curve(y_test,model_dt.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr,tpr,label='Random Forest Classifier (area = %0.2f)' % logit_roc_auc)
plt.plot([0,1],[0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```




In [ ]:

```
In [ ]: # initialize model : finding optimzal k value with SMOTE data
model2 = RandomForestClassifier(random_state=7)

cv_scores = []
k_values=range(2, 20)
# selecting best value of k
for k in k_values:
    kf=StratifiedKFold(n_splits=k,shuffle=True,random_state=7)
    scores=cross_val_score(model2, X_train_res, y_train_res,cv=kf)
    cv_scores.append(np.mean(scores))

# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, cv_scores, marker='o')
plt.xlabel('Number of Folds (k)')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Cross-Validated Accuracy for Different Values of k')
plt.grid()
plt.show()

# Find the best value of k
best_k = k_values[np.argmax(cv_scores)]
print(f"Best value of k: {best_k} with accuracy: {max(cv_scores):.4f}")
```



Best value of k: 13 with accuracy: 0.9373

```
In [ ]: # initialize model : training base model with SMOTE data
kfold = StratifiedKFold(n_splits=best_k)

rf_clf2 = RandomForestClassifier(random_state=7)
cv_results_rf2 = cross_validate(rf_clf, X_train_res, y_train_res, cv=kfold, scoring='accuracy')

print(f"K-Fold Accuracy Mean: \n Train: {cv_results_rf2['train_score'].mean()*100}%, \n Validation: {cv_results_rf2['cv_score'].mean()*100}%")
print(f"K-Fold Accuracy Std: \n Train: {cv_results_rf2['train_score'].std()*100}%, \n Validation: {cv_results_rf2['cv_score'].std()*100}%")
```

K-Fold Accuracy Mean:

Train: 100.00

Validation: 93.54

K-Fold Accuracy Std:

Train: 0.00,


Validation: 2.33

```
In [ ]: # selecting best estimators
params = {
    'n_estimators' : [100,200,300,400],
    'max_depth' : [3,5,10],
    'criterion' : ['gini', 'entropy'],
    'bootstrap' : [True, False],
    'max_features' : [8,9,10]
}

grid_rf2 = GridSearchCV(estimator = RandomForestClassifier(random_state=7),
                        param_grid = params,
                        scoring = 'accuracy',
                        cv = kfold,
                        n_jobs=-1)

grid_rf2.fit(X_train, y_train)
```

Out[ ]:




```
In [ ]: print("Best params: ", grid_rf2.best_params_)
        print("Best score: ", grid_rf2.best_score_)
```

```
Best params: {'bootstrap': True, 'criterion': 'entropy', 'max_depth': 10, 'max_f
eatures': 8, 'n_estimators': 300}
Best score: 0.9238331792147838
```

```
In [ ]: model_rf2 = grid_rf2.best_estimator_
        y_pred = model_rf2.predict(X_test)

        cm = confusion_matrix(y_test, y_pred)
        disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                                       display_labels=model_rf2.classes_)

        disp.plot()
        plt.show()
```




```
In [ ]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.90	0.90	153
1	0.95	0.95	0.95	324
accuracy			0.94	477
macro avg	0.93	0.93	0.93	477
weighted avg	0.94	0.94	0.94	477

```
In [ ]: rf_model2=classification_report(y_test, y_pred,output_dict=True)
```


```
In [ ]: importances = model_rf2.feature_importances_
indices = np.argsort(importances[::-1])
names = [X.columns[i] for i in indices]

plt.figure(figsize=(15, 7))
plt.title("Feature Importance")
plt.bar(range(X_train.shape[1]), importances[indices])
plt.xticks(range(X_train.shape[1]), names, rotation=45)
plt.show()
```



```
In [ ]: logit_roc_auc=roc_auc_score(y_test,y_pred)
fpr, tpr, thresholds=roc_curve(y_test,model_rf2.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr,tpr,label='Random Forest Classifier (area = %0.2f)' % logit_roc_auc)
plt.plot([0,1],[0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```





model with class\_weight and model with SMOTE performs similar


## Model 2 : GradientBoostingClassifier

```
In [ ]: # initialize model : finding best value for k with GBDT
model3 = GradientBoostingClassifier(random_state=7)

cv_scores = []
k_values=range(2, 20)
# selecting best value of k
for k in k_values:
    kf=StratifiedKFold(n_splits=k,shuffle=True,random_state=7)
    scores=cross_val_score(model3, X_train_res, y_train_res,cv=kf)
    cv_scores.append(np.mean(scores))

# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, cv_scores, marker='o')
plt.xlabel('Number of Folds (k)')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Cross-Validated Accuracy for Different Values of k')
plt.grid()
plt.show()

# Find the best value of k
best_k = k_values[np.argmax(cv_scores)]
print(f"Best value of k: {best_k} with accuracy: {max(cv_scores):.4f}")
```



Best value of k: 18 with accuracy: 0.9381

```
In [ ]: # training the base model with optimal k
kfold = StratifiedKFold(n_splits=best_k, shuffle=True, random_state=7)

gbdt_clf = GradientBoostingClassifier(random_state=7)
cv_results_rf = cross_validate(gbdt_clf, X_train_res, y_train_res, cv=kfold, scoring='accuracy')

print(f"K-Fold Accuracy Mean: \n Train: {cv_results_rf['train_score'].mean()*100:.2f}% \n Validation: {cv_results_rf['cv_mean_score'].mean()*100:.2f}%")
print(f"K-Fold Accuracy Std: \n Train: {cv_results_rf['train_score'].std()*100:.2f}% \n Validation: {cv_results_rf['cv_std_score'].std()*100:.2f}%")
```

K-Fold Accuracy Mean:

Train: 95.76

Validation: 93.65

K-Fold Accuracy Std:

Train: 0.17

Validation: 2.38

```
In [ ]: # selecting best params for the estimator
params = {
    'n_estimators': [100, 200, 300, 400],
    'max_depth': [3, 5, 10],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_features': [8, 9, 10]
}

grid_gbdt = GridSearchCV(estimator=GradientBoostingClassifier(random_state=7),
                          param_grid=params,
                          scoring='accuracy',
                          cv=kfold,
                          n_jobs=-1)

grid_gbdt.fit(X_train, y_train)
```

Out[ ]: **GridSearchCV**

- ▶ **best\_estimator\_: GradientBoostingClassifier**
  - ▶ GradientBoostingClassifier

```
In [ ]: print("Best params: ", grid_gbdtd.best_params_)
        print("Best score: ", grid_gbdtd.best_score_)
```


Best params: {'learning\_rate': 0.1, 'max\_depth': 5, 'max\_features': 9, 'n\_estimators': 200}

Best score: 0.9233003893381251

```
In [ ]: model_gbdtd = grid_gbdtd.best_estimator_
        y_pred = model_gbdtd.predict(X_test)

        cm = confusion_matrix(y_test, y_pred)
        disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                                       display_labels=model_rf2.classes_)

        disp.plot()
        plt.show()
```




```
In [ ]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.92	0.90	153
1	0.96	0.94	0.95	324
accuracy			0.94	477
macro avg	0.92	0.93	0.93	477
weighted avg	0.94	0.94	0.94	477


```
In [ ]: gbd_model=classification_report(y_test, y_pred,output_dict=True)
```

```
In [ ]: importances = model_gbd.feature_importances_
indices = np.argsort(importances)[::-1]
names = [X.columns[i] for i in indices]

plt.figure(figsize=(15, 7))
plt.title("Feature Importance")
plt.bar(range(X_train.shape[1]), importances[indices])
plt.xticks(range(X_train.shape[1]), names, rotation=45)
plt.show()
```



```
In [ ]: logit_roc_auc=roc_auc_score(y_test,y_pred)
fpr, tpr, thresholds=roc_curve(y_test,model_rf2.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr,tpr,label='Random Forest Classifier (area = %0.2f)' % logit_roc_auc)
plt.plot([0,1],[0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```




### Model 3 : XGBoost

```
In [ ]: # initialize model : finding best value for k with GBDT
model4 = XGBClassifier(random_state=7)

cv_scores = []
k_values=range(2, 20)
# selecting best value of k
for k in k_values:
    kf=StratifiedKFold(n_splits=k,shuffle=True,random_state=7)
    scores=cross_val_score(model4, X_train_res, y_train_res,cv=kf)
    cv_scores.append(np.mean(scores))

# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, cv_scores, marker='o')
plt.xlabel('Number of Folds (k)')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Cross-Validated Accuracy for Different Values of k')
plt.grid()
plt.show()

# Find the best value of k
best_k = k_values[np.argmax(cv_scores)]
print(f"Best value of k: {best_k} with accuracy: {max(cv_scores):.4f}")
```



Best value of k: 17 with accuracy: 0.9443

```
In [ ]: # training the base model with optimal k
kfold = StratifiedKFold(n_splits=best_k, shuffle=True, random_state=7)

xgb_clf = XGBClassifier(random_state=7)
cv_results_rf = cross_validate(xgb_clf, X_train_res, y_train_res, cv=kfold, scoring='accuracy')

print(f"K-Fold Accuracy Mean: \n Train: {cv_results_rf['train_score'].mean()*100:.2f}%")
print(f"K-Fold Accuracy Std: \n Train: {cv_results_rf['train_score'].std()*100:.2f}%")
```

K-Fold Accuracy Mean:

Train: 100.00

Validation: 94.43

K-Fold Accuracy Std:

Train: 0.00

Validation: 2.20

```
In [ ]: # selecting best params for the estimator
params = {
    'n_estimators': [100, 200, 300, 400],
    'max_depth': [3, 5, 10],
    'learning_rate': [0.01, 0.1, 0.2],
    'gamma': [0, 0.1, 0.2],
    'reg_alpha': [0, 0.01, 0.1],
    'reg_lambda': [1, 1.5, 2]
}

grid_xgb = GridSearchCV(estimator=XGBClassifier(random_state=7),
                        param_grid=params,
                        scoring='accuracy',
                        cv=kfold,
                        n_jobs=-1)

grid_xgb.fit(X_train, y_train)
```

Out[ ]:

```

GridSearchCV
  best_estimator_: XGBClassifier
    XGBClassifier

```

```

In [ ]: print("Best params: ", grid_xgb.best_params_)
        print("Best score: ", grid_xgb.best_score_)

```

Best params: {'gamma': 0, 'learning\_rate': 0.2, 'max\_depth': 10, 'n\_estimators': 200, 'reg\_alpha': 0, 'reg\_lambda': 1.5}  
 Best score: 0.928046218487395


```

In [ ]: model_xgb = grid_xgb.best_estimator_
        y_pred = model_xgb.predict(X_test)

        cm = confusion_matrix(y_test, y_pred)
        disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                                       display_labels=model_xgb.classes_)

        disp.plot()
        plt.show()

```



```

In [ ]: print(classification_report(y_test, y_pred))


```

	precision	recall	f1-score	support
0	0.88	0.92	0.89	153
1	0.96	0.94	0.95	324
accuracy			0.93	477
macro avg	0.92	0.93	0.92	477
weighted avg	0.93	0.93	0.93	477

```
In [ ]: xbg_model=classification_report(y_test, y_pred,output_dict=True)
```


```
In [ ]: importances = model_xgb.feature_importances_
indices = np.argsort(importances[::-1])
names = [X.columns[i] for i in indices]

plt.figure(figsize=(15, 7))
plt.title("Feature Importance")
plt.bar(range(X_train.shape[1]), importances[indices])
plt.xticks(range(X_train.shape[1]), names, rotation=45)
plt.show()
```



```
In [ ]: logit_roc_auc=roc_auc_score(y_test,y_pred)
fpr, tpr, thresholds=roc_curve(y_test,model_rf2.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr,tpr,label='Random Forest Classifier (area = %0.2f)' % logit_roc_auc)
plt.plot([0,1],[0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```





### Model 4 : LightGBM

```
In [ ]: # initialize model : finding best value for k with GBDT
model5 = LGBMClassifier(random_state=7)

cv_scores = []
k_values=range(2, 20)
# selecting best value of k
for k in k_values:
    kf=StratifiedKFold(n_splits=k,shuffle=True,random_state=7)
    scores=cross_val_score(model5, X_train_res, y_train_res,cv=kf)
    cv_scores.append(np.mean(scores))

# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, cv_scores, marker='o')
plt.xlabel('Number of Folds (k)')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Cross-Validated Accuracy for Different Values of k')
plt.grid()
plt.show()

# Find the best value of k
best_k = k_values[np.argmax(cv_scores)]
print(f"Best value of k: {best_k} with accuracy: {max(cv_scores):.4f}")
```

[LightGBM] [Info] Number of positive: 646, number of negative: 646  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000929 seconds.  
You can set `force\_row\_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force\_col\_wise=true`.  
[LightGBM] [Info] Total Bins 2056  
[LightGBM] [Info] Number of data points in the train set: 1292, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf  
[LightGBM] [Info] Number of positive: 646, number of negative: 646  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000155 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2057  
[LightGBM] [Info] Number of data points in the train set: 1292, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 861, number of negative: 861  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000329 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2430  
[LightGBM] [Info] Number of data points in the train set: 1722, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 862, number of negative: 861  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000223 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2433  
[LightGBM] [Info] Number of data points in the train set: 1723, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500290 -> initscore=0.001161  
[LightGBM] [Info] Start training from score 0.001161  
[LightGBM] [Info] Number of positive: 861, number of negative: 862  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000190 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2422  
[LightGBM] [Info] Number of data points in the train set: 1723, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499710 -> initscore=-0.001161  
[LightGBM] [Info] Start training from score -0.001161  
[LightGBM] [Info] Number of positive: 969, number of negative: 969  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000199 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2512  
[LightGBM] [Info] Number of data points in the train set: 1938, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 969, number of negative: 969  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000168 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2454  
[LightGBM] [Info] Number of data points in the train set: 1938, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 969, number of negative: 969  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000197 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2521  
[LightGBM] [Info] Number of data points in the train set: 1938, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 969, number of negative: 969  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000240 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2504  
[LightGBM] [Info] Number of data points in the train set: 1938, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1034, number of negative: 1033  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000214 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2521  
[LightGBM] [Info] Number of data points in the train set: 2067, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500242 -> initscore=0.000968  
[LightGBM] [Info] Start training from score 0.000968  
[LightGBM] [Info] Number of positive: 1034, number of negative: 1033  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000206 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2531  
[LightGBM] [Info] Number of data points in the train set: 2067, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500242 -> initscore=0.000968  
[LightGBM] [Info] Start training from score 0.000968  
[LightGBM] [Info] Number of positive: 1033, number of negative: 1034  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000233 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2533  
[LightGBM] [Info] Number of data points in the train set: 2067, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499758 -> initscore=-0.000968  
[LightGBM] [Info] Start training from score -0.000968  
[LightGBM] [Info] Number of positive: 1033, number of negative: 1034  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000237 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2514  
[LightGBM] [Info] Number of data points in the train set: 2067, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499758 -> initscore=-0.000968  
[LightGBM] [Info] Start training from score -0.000968  
[LightGBM] [Info] Number of positive: 1034, number of negative: 1034  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000192 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2527  
[LightGBM] [Info] Number of data points in the train set: 2068, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

```
[LightGBM] [Info] Number of positive: 1077, number of negative: 1076
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000204 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2530
[LightGBM] [Info] Number of data points in the train set: 2153, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500232 -> initscore=0.000929
[LightGBM] [Info] Start training from score 0.000929
[LightGBM] [Info] Number of positive: 1077, number of negative: 1076
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000208 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2534
[LightGBM] [Info] Number of data points in the train set: 2153, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500232 -> initscore=0.000929
[LightGBM] [Info] Start training from score 0.000929
[LightGBM] [Info] Number of positive: 1076, number of negative: 1077
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000188 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2530
[LightGBM] [Info] Number of data points in the train set: 2153, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499768 -> initscore=-0.000929
[LightGBM] [Info] Start training from score -0.000929
[LightGBM] [Info] Number of positive: 1076, number of negative: 1077
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000180 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2543
[LightGBM] [Info] Number of data points in the train set: 2153, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499768 -> initscore=-0.000929
[LightGBM] [Info] Start training from score -0.000929
[LightGBM] [Info] Number of positive: 1077, number of negative: 1077
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000192 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2526
[LightGBM] [Info] Number of data points in the train set: 2154, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1077, number of negative: 1077
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing
was 0.000237 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2536
[LightGBM] [Info] Number of data points in the train set: 2154, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1107, number of negative: 1107
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000235 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2538
[LightGBM] [Info] Number of data points in the train set: 2214, number of used fe
atures: 16
```

```

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1108, number of negative: 1107
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000207 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2545
[LightGBM] [Info] Number of data points in the train set: 2215, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500226 -> initscore=0.000903
[LightGBM] [Info] Start training from score 0.000903
[LightGBM] [Info] Number of positive: 1108, number of negative: 1107
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000176 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2538
[LightGBM] [Info] Number of data points in the train set: 2215, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500226 -> initscore=0.000903
[LightGBM] [Info] Start training from score 0.000903
[LightGBM] [Info] Number of positive: 1108, number of negative: 1107
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000187 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2550
[LightGBM] [Info] Number of data points in the train set: 2215, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500226 -> initscore=0.000903
[LightGBM] [Info] Start training from score 0.000903
[LightGBM] [Info] Number of positive: 1107, number of negative: 1108
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000175 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2539
[LightGBM] [Info] Number of data points in the train set: 2215, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499774 -> initscore=-0.000903
[LightGBM] [Info] Start training from score -0.000903
[LightGBM] [Info] Number of positive: 1107, number of negative: 1108
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000197 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2533
[LightGBM] [Info] Number of data points in the train set: 2215, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499774 -> initscore=-0.000903
[LightGBM] [Info] Start training from score -0.000903
[LightGBM] [Info] Number of positive: 1107, number of negative: 1108
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000220 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2544
[LightGBM] [Info] Number of data points in the train set: 2215, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499774 -> initscore=-0.000903
[LightGBM] [Info] Start training from score -0.000903
[LightGBM] [Info] Number of positive: 1131, number of negative: 1130
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000197 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2542

```

```
[LightGBM] [Info] Number of data points in the train set: 2261, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500221 -> initscore=0.000885
[LightGBM] [Info] Start training from score 0.000885
[LightGBM] [Info] Number of positive: 1131, number of negative: 1130
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000203 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2549
[LightGBM] [Info] Number of data points in the train set: 2261, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500221 -> initscore=0.000885
[LightGBM] [Info] Start training from score 0.000885
[LightGBM] [Info] Number of positive: 1131, number of negative: 1130
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000225 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2547
[LightGBM] [Info] Number of data points in the train set: 2261, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500221 -> initscore=0.000885
[LightGBM] [Info] Start training from score 0.000885
[LightGBM] [Info] Number of positive: 1131, number of negative: 1130
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000200 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2547
[LightGBM] [Info] Number of data points in the train set: 2261, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500221 -> initscore=0.000885
[LightGBM] [Info] Start training from score 0.000885
[LightGBM] [Info] Number of positive: 1130, number of negative: 1131
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000277 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2261, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499779 -> initscore=-0.000885
[LightGBM] [Info] Start training from score -0.000885
[LightGBM] [Info] Number of positive: 1130, number of negative: 1131
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000196 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2540
[LightGBM] [Info] Number of data points in the train set: 2261, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499779 -> initscore=-0.000885
[LightGBM] [Info] Start training from score -0.000885
[LightGBM] [Info] Number of positive: 1130, number of negative: 1131
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000225 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2530
[LightGBM] [Info] Number of data points in the train set: 2261, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499779 -> initscore=-0.000885
[LightGBM] [Info] Start training from score -0.000885
[LightGBM] [Info] Number of positive: 1130, number of negative: 1131
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
```

was 0.000219 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2550

[LightGBM] [Info] Number of data points in the train set: 2261, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499779 -> initscore=-0.000885

[LightGBM] [Info] Start training from score -0.000885

[LightGBM] [Info] Number of positive: 1148, number of negative: 1148

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000150 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2544

[LightGBM] [Info] Number of data points in the train set: 2296, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 1149, number of negative: 1148

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000215 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2545

[LightGBM] [Info] Number of data points in the train set: 2297, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500218 -> initscore=0.000871

[LightGBM] [Info] Start training from score 0.000871

[LightGBM] [Info] Number of positive: 1149, number of negative: 1148

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000189 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2548

[LightGBM] [Info] Number of data points in the train set: 2297, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500218 -> initscore=0.000871

[LightGBM] [Info] Start training from score 0.000871

[LightGBM] [Info] Number of positive: 1149, number of negative: 1148

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000252 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2550

[LightGBM] [Info] Number of data points in the train set: 2297, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500218 -> initscore=0.000871

[LightGBM] [Info] Start training from score 0.000871

[LightGBM] [Info] Number of positive: 1149, number of negative: 1148

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000237 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2562

[LightGBM] [Info] Number of data points in the train set: 2297, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500218 -> initscore=0.000871

[LightGBM] [Info] Start training from score 0.000871

[LightGBM] [Info] Number of positive: 1148, number of negative: 1149

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000228 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2553

[LightGBM] [Info] Number of data points in the train set: 2297, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499782 -> initscore=-0.000871

[LightGBM] [Info] Start training from score -0.000871

```
[LightGBM] [Info] Number of positive: 1148, number of negative: 1149
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000226 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2540
[LightGBM] [Info] Number of data points in the train set: 2297, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499782 -> initscore=-0.000871
[LightGBM] [Info] Start training from score -0.000871
[LightGBM] [Info] Number of positive: 1148, number of negative: 1149
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000204 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2541
[LightGBM] [Info] Number of data points in the train set: 2297, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499782 -> initscore=-0.000871
[LightGBM] [Info] Start training from score -0.000871
[LightGBM] [Info] Number of positive: 1148, number of negative: 1149
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000226 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2552
[LightGBM] [Info] Number of data points in the train set: 2297, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499782 -> initscore=-0.000871
[LightGBM] [Info] Start training from score -0.000871
[LightGBM] [Info] Number of positive: 1163, number of negative: 1162
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000222 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2547
[LightGBM] [Info] Number of data points in the train set: 2325, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500215 -> initscore=0.000860
[LightGBM] [Info] Start training from score 0.000860
[LightGBM] [Info] Number of positive: 1163, number of negative: 1162
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000202 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2545
[LightGBM] [Info] Number of data points in the train set: 2325, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500215 -> initscore=0.000860
[LightGBM] [Info] Start training from score 0.000860
[LightGBM] [Info] Number of positive: 1162, number of negative: 1163
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000238 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2552
[LightGBM] [Info] Number of data points in the train set: 2325, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499785 -> initscore=-0.000860
[LightGBM] [Info] Start training from score -0.000860
[LightGBM] [Info] Number of positive: 1162, number of negative: 1163
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000202 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2544
[LightGBM] [Info] Number of data points in the train set: 2325, number of used fe
```



```

atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499785 -> initscore=-0.000860
[LightGBM] [Info] Start training from score -0.000860
[LightGBM] [Info] Number of positive: 1163, number of negative: 1163
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000218 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2549
[LightGBM] [Info] Number of data points in the train set: 2326, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1163, number of negative: 1163
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing
was 0.000243 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2560
[LightGBM] [Info] Number of data points in the train set: 2326, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1163, number of negative: 1163
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000217 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2553
[LightGBM] [Info] Number of data points in the train set: 2326, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1163, number of negative: 1163
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000181 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2542
[LightGBM] [Info] Number of data points in the train set: 2326, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1163, number of negative: 1163
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000219 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2326, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1163, number of negative: 1163
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000187 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2326, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1175, number of negative: 1174
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000201 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2550
[LightGBM] [Info] Number of data points in the train set: 2349, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500213 -> initscore=0.000851

```

```
[LightGBM] [Info] Start training from score 0.000851
[LightGBM] [Info] Number of positive: 1175, number of negative: 1174
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000198 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2552
[LightGBM] [Info] Number of data points in the train set: 2349, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500213 -> initscore=0.000851
[LightGBM] [Info] Start training from score 0.000851
[LightGBM] [Info] Number of positive: 1175, number of negative: 1174
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000203 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2556
[LightGBM] [Info] Number of data points in the train set: 2349, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500213 -> initscore=0.000851
[LightGBM] [Info] Start training from score 0.000851
[LightGBM] [Info] Number of positive: 1175, number of negative: 1174
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000197 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2553
[LightGBM] [Info] Number of data points in the train set: 2349, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500213 -> initscore=0.000851
[LightGBM] [Info] Start training from score 0.000851
[LightGBM] [Info] Number of positive: 1175, number of negative: 1174
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000190 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2349, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500213 -> initscore=0.000851
[LightGBM] [Info] Start training from score 0.000851
[LightGBM] [Info] Number of positive: 1174, number of negative: 1175
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000230 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2559
[LightGBM] [Info] Number of data points in the train set: 2349, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499787 -> initscore=-0.000851
[LightGBM] [Info] Start training from score -0.000851
[LightGBM] [Info] Number of positive: 1174, number of negative: 1175
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000229 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2559
[LightGBM] [Info] Number of data points in the train set: 2349, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499787 -> initscore=-0.000851
[LightGBM] [Info] Start training from score -0.000851
[LightGBM] [Info] Number of positive: 1174, number of negative: 1175
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000207 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2550
```

```
[LightGBM] [Info] Number of data points in the train set: 2349, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499787 -> initscore=-0.000851
[LightGBM] [Info] Start training from score -0.000851
[LightGBM] [Info] Number of positive: 1174, number of negative: 1175
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000205 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2545
[LightGBM] [Info] Number of data points in the train set: 2349, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499787 -> initscore=-0.000851
[LightGBM] [Info] Start training from score -0.000851
[LightGBM] [Info] Number of positive: 1174, number of negative: 1175
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000237 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2552
[LightGBM] [Info] Number of data points in the train set: 2349, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499787 -> initscore=-0.000851
[LightGBM] [Info] Start training from score -0.000851
[LightGBM] [Info] Number of positive: 1175, number of negative: 1175
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000222 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2558
[LightGBM] [Info] Number of data points in the train set: 2350, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1184, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000212 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2552
[LightGBM] [Info] Number of data points in the train set: 2368, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1184, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000184 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2368, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1184, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000236 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2557
[LightGBM] [Info] Number of data points in the train set: 2368, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1184, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000223 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2558
[LightGBM] [Info] Number of data points in the train set: 2368, number of used fe
```

```

atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1185, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000190 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2557
[LightGBM] [Info] Number of data points in the train set: 2369, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500211 -> initscore=0.000844
[LightGBM] [Info] Start training from score 0.000844
[LightGBM] [Info] Number of positive: 1185, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000243 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2556
[LightGBM] [Info] Number of data points in the train set: 2369, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500211 -> initscore=0.000844
[LightGBM] [Info] Start training from score 0.000844
[LightGBM] [Info] Number of positive: 1185, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000218 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2568
[LightGBM] [Info] Number of data points in the train set: 2369, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500211 -> initscore=0.000844
[LightGBM] [Info] Start training from score 0.000844
[LightGBM] [Info] Number of positive: 1185, number of negative: 1184
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000250 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2369, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500211 -> initscore=0.000844
[LightGBM] [Info] Start training from score 0.000844
[LightGBM] [Info] Number of positive: 1184, number of negative: 1185
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000245 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2550
[LightGBM] [Info] Number of data points in the train set: 2369, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499789 -> initscore=-0.000844
[LightGBM] [Info] Start training from score -0.000844
[LightGBM] [Info] Number of positive: 1184, number of negative: 1185
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000235 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2546
[LightGBM] [Info] Number of data points in the train set: 2369, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499789 -> initscore=-0.000844
[LightGBM] [Info] Start training from score -0.000844
[LightGBM] [Info] Number of positive: 1184, number of negative: 1185
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000195 seconds.
You can set `force_col_wise=true` to remove the overhead.

```

```
[LightGBM] [Info] Total Bins 2554
[LightGBM] [Info] Number of data points in the train set: 2369, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499789 -> initscore=-0.000844
[LightGBM] [Info] Start training from score -0.000844
[LightGBM] [Info] Number of positive: 1184, number of negative: 1185
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000248 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2561
[LightGBM] [Info] Number of data points in the train set: 2369, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499789 -> initscore=-0.000844
[LightGBM] [Info] Start training from score -0.000844
[LightGBM] [Info] Number of positive: 1193, number of negative: 1192
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000189 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2556
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500210 -> initscore=0.000839
[LightGBM] [Info] Start training from score 0.000839
[LightGBM] [Info] Number of positive: 1193, number of negative: 1192
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000206 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2558
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500210 -> initscore=0.000839
[LightGBM] [Info] Start training from score 0.000839
[LightGBM] [Info] Number of positive: 1193, number of negative: 1192
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000241 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2559
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500210 -> initscore=0.000839
[LightGBM] [Info] Start training from score 0.000839
[LightGBM] [Info] Number of positive: 1193, number of negative: 1192
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000229 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2554
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500210 -> initscore=0.000839
[LightGBM] [Info] Start training from score 0.000839
[LightGBM] [Info] Number of positive: 1193, number of negative: 1192
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000201 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2551
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500210 -> initscore=0.000839
[LightGBM] [Info] Start training from score 0.000839
[LightGBM] [Info] Number of positive: 1192, number of negative: 1193
```

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000189 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2560  
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499790 -> initscore=-0.000839  
[LightGBM] [Info] Start training from score -0.000839  
[LightGBM] [Info] Number of positive: 1192, number of negative: 1193  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000238 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2562  
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499790 -> initscore=-0.000839  
[LightGBM] [Info] Start training from score -0.000839  
[LightGBM] [Info] Number of positive: 1192, number of negative: 1193  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000263 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2565  
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499790 -> initscore=-0.000839  
[LightGBM] [Info] Start training from score -0.000839  
[LightGBM] [Info] Number of positive: 1192, number of negative: 1193  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000190 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2556  
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499790 -> initscore=-0.000839  
[LightGBM] [Info] Start training from score -0.000839  
[LightGBM] [Info] Number of positive: 1192, number of negative: 1193  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000214 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2553  
[LightGBM] [Info] Number of data points in the train set: 2385, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499790 -> initscore=-0.000839  
[LightGBM] [Info] Start training from score -0.000839  
[LightGBM] [Info] Number of positive: 1193, number of negative: 1193  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000262 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2544  
[LightGBM] [Info] Number of data points in the train set: 2386, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1193, number of negative: 1193  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000217 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2720  
[LightGBM] [Info] Number of data points in the train set: 2386, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 1193, number of negative: 1193  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000296 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2564  
[LightGBM] [Info] Number of data points in the train set: 2386, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1200, number of negative: 1199  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000200 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2558  
[LightGBM] [Info] Number of data points in the train set: 2399, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500208 -> initscore=0.000834  
[LightGBM] [Info] Start training from score 0.000834  
[LightGBM] [Info] Number of positive: 1200, number of negative: 1199  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000242 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2560  
[LightGBM] [Info] Number of data points in the train set: 2399, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500208 -> initscore=0.000834  
[LightGBM] [Info] Start training from score 0.000834  
[LightGBM] [Info] Number of positive: 1200, number of negative: 1199  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000210 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2557  
[LightGBM] [Info] Number of data points in the train set: 2399, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500208 -> initscore=0.000834  
[LightGBM] [Info] Start training from score 0.000834  
[LightGBM] [Info] Number of positive: 1200, number of negative: 1199  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000209 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2563  
[LightGBM] [Info] Number of data points in the train set: 2399, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500208 -> initscore=0.000834  
[LightGBM] [Info] Start training from score 0.000834  
[LightGBM] [Info] Number of positive: 1199, number of negative: 1200  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000233 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2558  
[LightGBM] [Info] Number of data points in the train set: 2399, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499792 -> initscore=-0.000834  
[LightGBM] [Info] Start training from score -0.000834  
[LightGBM] [Info] Number of positive: 1199, number of negative: 1200  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000224 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2558  
[LightGBM] [Info] Number of data points in the train set: 2399, number of used features: 16

```
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499792 -> initscore=-0.000834
[LightGBM] [Info] Start training from score -0.000834
[LightGBM] [Info] Number of positive: 1199, number of negative: 1200
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000198 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2399, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499792 -> initscore=-0.000834
[LightGBM] [Info] Start training from score -0.000834
[LightGBM] [Info] Number of positive: 1199, number of negative: 1200
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000235 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2570
[LightGBM] [Info] Number of data points in the train set: 2399, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499792 -> initscore=-0.000834
[LightGBM] [Info] Start training from score -0.000834
[LightGBM] [Info] Number of positive: 1200, number of negative: 1200
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000228 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2562
[LightGBM] [Info] Number of data points in the train set: 2400, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1200, number of negative: 1200
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000179 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2400, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1200, number of negative: 1200
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000210 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2551
[LightGBM] [Info] Number of data points in the train set: 2400, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1200, number of negative: 1200
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000201 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2557
[LightGBM] [Info] Number of data points in the train set: 2400, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1200, number of negative: 1200
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000278 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2721
[LightGBM] [Info] Number of data points in the train set: 2400, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
```



[LightGBM] [Info] Number of positive: 1200, number of negative: 1200  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000246 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2567  
[LightGBM] [Info] Number of data points in the train set: 2400, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1205  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000218 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2560  
[LightGBM] [Info] Number of data points in the train set: 2411, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500207 -> initscore=0.000830  
[LightGBM] [Info] Start training from score 0.000830  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1205  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000194 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2563  
[LightGBM] [Info] Number of data points in the train set: 2411, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500207 -> initscore=0.000830  
[LightGBM] [Info] Start training from score 0.000830  
[LightGBM] [Info] Number of positive: 1205, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000171 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2554  
[LightGBM] [Info] Number of data points in the train set: 2411, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499793 -> initscore=-0.000830  
[LightGBM] [Info] Start training from score -0.000830  
[LightGBM] [Info] Number of positive: 1205, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000232 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2564  
[LightGBM] [Info] Number of data points in the train set: 2411, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499793 -> initscore=-0.000830  
[LightGBM] [Info] Start training from score -0.000830  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000241 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2561  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000193 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2558  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000283 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2562  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000206 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2565  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000231 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2565  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000236 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2560  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000192 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2557  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000208 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2552  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000240 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2564  
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000266 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

```
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1206, number of negative: 1206
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000237 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2567
[LightGBM] [Info] Number of data points in the train set: 2412, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1211, number of negative: 1211
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000211 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1211, number of negative: 1211
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000194 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2561
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1211, number of negative: 1211
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000222 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2558
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1211, number of negative: 1211
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000199 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2566
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1211, number of negative: 1211
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000193 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1211, number of negative: 1211
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000264 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2561
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
```

[LightGBM] [Info] Number of positive: 1211, number of negative: 1211  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000223 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2563  
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1211, number of negative: 1211  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000160 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2565  
[LightGBM] [Info] Number of data points in the train set: 2422, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1212, number of negative: 1211  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000237 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2572  
[LightGBM] [Info] Number of data points in the train set: 2423, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500206 -> initscore=0.000825  
[LightGBM] [Info] Start training from score 0.000825  
[LightGBM] [Info] Number of positive: 1212, number of negative: 1211  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000248 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2563  
[LightGBM] [Info] Number of data points in the train set: 2423, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500206 -> initscore=0.000825  
[LightGBM] [Info] Start training from score 0.000825  
[LightGBM] [Info] Number of positive: 1212, number of negative: 1211  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000215 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2560  
[LightGBM] [Info] Number of data points in the train set: 2423, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500206 -> initscore=0.000825  
[LightGBM] [Info] Start training from score 0.000825  
[LightGBM] [Info] Number of positive: 1212, number of negative: 1211  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000184 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2556  
[LightGBM] [Info] Number of data points in the train set: 2423, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500206 -> initscore=0.000825  
[LightGBM] [Info] Start training from score 0.000825  
[LightGBM] [Info] Number of positive: 1211, number of negative: 1212  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000208 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2552  
[LightGBM] [Info] Number of data points in the train set: 2423, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499794 -> initscore=-0.000825

```
[LightGBM] [Info] Start training from score -0.000825
[LightGBM] [Info] Number of positive: 1211, number of negative: 1212
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000209 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2566
[LightGBM] [Info] Number of data points in the train set: 2423, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499794 -> initscore=-0.000825
[LightGBM] [Info] Start training from score -0.000825
[LightGBM] [Info] Number of positive: 1211, number of negative: 1212
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000236 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2560
[LightGBM] [Info] Number of data points in the train set: 2423, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499794 -> initscore=-0.000825
[LightGBM] [Info] Start training from score -0.000825
[LightGBM] [Info] Number of positive: 1211, number of negative: 1212
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000211 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2571
[LightGBM] [Info] Number of data points in the train set: 2423, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499794 -> initscore=-0.000825
[LightGBM] [Info] Start training from score -0.000825
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000215 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2565
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000217 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2558
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000202 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000187 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2566
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
```

[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000218 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2564  
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000227 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2564  
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000240 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2568  
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000232 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2561  
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000272 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2568  
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000185 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2569  
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000232 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 2565  
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16  
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000233 seconds.  
You can set `force\_col\_wise=true` to remove the overhead.

```
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000216 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000236 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000201 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2566
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000280 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2562
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000204 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2571
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000233 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2440, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000189 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2562
[LightGBM] [Info] Number of data points in the train set: 2440, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
```

```
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000279 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2560
[LightGBM] [Info] Number of data points in the train set: 2440, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000204 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2566
[LightGBM] [Info] Number of data points in the train set: 2440, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000203 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2569
[LightGBM] [Info] Number of data points in the train set: 2440, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000225 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2440, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000203 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2440, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000239 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2565
[LightGBM] [Info] Number of data points in the train set: 2440, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000241 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2567
[LightGBM] [Info] Number of data points in the train set: 2440, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1220, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000215 seconds.
You can set `force_col_wise=true` to remove the overhead.
```



```
[LightGBM] [Info] Total Bins 2572
[LightGBM] [Info] Number of data points in the train set: 2440, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1221, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000230 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2568
[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500205 -> initscore=0.000819
[LightGBM] [Info] Start training from score 0.000819
[LightGBM] [Info] Number of positive: 1221, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000242 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500205 -> initscore=0.000819
[LightGBM] [Info] Start training from score 0.000819
[LightGBM] [Info] Number of positive: 1221, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000249 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2560
[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500205 -> initscore=0.000819
[LightGBM] [Info] Start training from score 0.000819
[LightGBM] [Info] Number of positive: 1221, number of negative: 1220
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000223 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500205 -> initscore=0.000819
[LightGBM] [Info] Start training from score 0.000819
[LightGBM] [Info] Number of positive: 1220, number of negative: 1221
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000324 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2560
[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499795 -> initscore=-0.000819
[LightGBM] [Info] Start training from score -0.000819
[LightGBM] [Info] Number of positive: 1220, number of negative: 1221
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000284 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2565
[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499795 -> initscore=-0.000819
[LightGBM] [Info] Start training from score -0.000819
[LightGBM] [Info] Number of positive: 1220, number of negative: 1221
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
```

was 0.000235 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2562

[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499795 -> initscore=-0.000819

[LightGBM] [Info] Start training from score -0.000819

[LightGBM] [Info] Number of positive: 1220, number of negative: 1221

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000214 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2571

[LightGBM] [Info] Number of data points in the train set: 2441, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499795 -> initscore=-0.000819

[LightGBM] [Info] Start training from score -0.000819

[LightGBM] [Info] Number of positive: 1224, number of negative: 1224

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000292 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2563

[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 1224, number of negative: 1224

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000227 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2566

[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 1224, number of negative: 1224

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000229 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2564

[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 1224, number of negative: 1224

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000227 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2566

[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[LightGBM] [Info] Number of positive: 1224, number of negative: 1224

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000189 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 2570

[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000


[LightGBM] [Info] Number of positive: 1224, number of negative: 1224

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000273 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

```
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000196 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2567
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000255 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2565
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000206 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2567
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000272 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2570
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000215 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2572
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000260 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000249 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2448, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
```

```
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000196 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2558
[LightGBM] [Info] Number of data points in the train set: 2448, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000229 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2557
[LightGBM] [Info] Number of data points in the train set: 2448, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000216 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2448, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing
was 0.000830 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2448, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000315 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2448, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1224, number of negative: 1224
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000356 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2571
[LightGBM] [Info] Number of data points in the train set: 2448, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
```



Best value of k: 17 with accuracy: 0.9478

```
In [ ]: # training the base model with optimal k
kfolds = StratifiedKFold(n_splits=best_k, shuffle=True, random_state=7)

lgbm_clf = LGBMClassifier(random_state=7)
cv_results_rf = cross_validate(lgbm_clf, X_train_res, y_train_res, cv=kfolds, scoring='accuracy')

print(f"K-Fold Accuracy Mean: \n Train: {cv_results_rf['train_score'].mean()*100:.2f}%")
print(f"K-Fold Accuracy Std: \n Train: {cv_results_rf['train_score'].std()*100:.2f}%")
```

```
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000232 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2565
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000237 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2558
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000163 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000196 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2566
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000232 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000218 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2564
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000240 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2568
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000205 seconds.
You can set `force_col_wise=true` to remove the overhead.
```

```
[LightGBM] [Info] Total Bins 2561
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000229 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2568
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000236 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2569
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000214 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2565
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000345 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2563
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000224 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000180 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2555
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000220 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2566
[LightGBM] [Info] Number of data points in the train set: 2432, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
```

```
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000243 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2562
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of positive: 1216, number of negative: 1216
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing
was 0.000389 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 2571
[LightGBM] [Info] Number of data points in the train set: 2432, number of used fe
atures: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
K-Fold Accuracy Mean:
  Train: 100.00
  Validation: 94.78
K-Fold Accuracy Std:
  Train: 0.00
  Validation: 2.10
```

```
In [ ]: # selecting best params for the estimator
params = {
    'n_estimators': [100, 200, 300, 400],
    'max_depth': [3, 5, 10],
    'learning_rate': [0.01, 0.1, 0.2],
    'reg_alpha': [0, 0.01, 0.1],
    'reg_lambda': [1, 1.5, 2]
}

grid_lgbm = GridSearchCV(estimator=LGBMClassifier(random_state=7),
                        param_grid=params,
                        scoring='accuracy',
                        cv=kfold,
                        n_jobs=-1)

grid_lgbm.fit(X_train, y_train)
```



```
[LightGBM] [Warning] Accuracy may be bad since you didn't explicitly set num_leaves OR 2^max_depth > num_leaves. (num_leaves=31).
[LightGBM] [Warning] Accuracy may be bad since you didn't explicitly set num_leaves OR 2^max_depth > num_leaves. (num_leaves=31).
[LightGBM] [Info] Number of positive: 1292, number of negative: 612
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000218 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1108
[LightGBM] [Info] Number of data points in the train set: 1904, number of used features: 16
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.678571 -> initscore=0.747214
[LightGBM] [Info] Start training from score 0.747214
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

```
Out[ ]:  ▶      GridSearchCV
          ▶ best_estimator_: LGBMClassifier
            ▶ LGBMClassifier
```

```
In [ ]: print("Best params: ", grid_lgbm.best_params_)
        print("Best score: ", grid_lgbm.best_score_)
```


```
Best params: {'learning_rate': 0.1, 'max_depth': 10, 'n_estimators': 100, 'reg_alpha': 0, 'reg_lambda': 1.5}
Best score: 0.9243697478991597
```

```
In [ ]: model_lgbm = grid_lgbm.best_estimator_
        y_pred = model_lgbm.predict(X_test)

        cm = confusion_matrix(y_test, y_pred)
        disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                                       display_labels=model_lgbm.classes_)

        disp.plot()
        plt.show()
```

```
[LightGBM] [Warning] Accuracy may be bad since you didn't explicitly set num_leaves OR 2^max_depth > num_leaves. (num_leaves=31).
```




```
In [ ]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.92	0.90	153
1	0.96	0.94	0.95	324
accuracy			0.94	477
macro avg	0.92	0.93	0.93	477
weighted avg	0.94	0.94	0.94	477


```
In [ ]: lgbm_model=classification_report(y_test, y_pred,output_dict=True)
```

```
In [ ]: importances = model_lgbm.feature_importances_
indices = np.argsort(importances)[::-1]
names = [X.columns[i] for i in indices]

plt.figure(figsize=(15, 7))
plt.title("Feature Importance")
plt.bar(range(X_train.shape[1]), importances[indices])
plt.xticks(range(X_train.shape[1]), names, rotation=45)
plt.show()
```



```
In [ ]: logit_roc_auc=roc_auc_score(y_test,y_pred)
fpr, tpr, thresholds=roc_curve(y_test,model_rf2.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr,tpr,label='Random Forest Classifier (area = %0.2f)' % logit_roc_auc)
plt.plot([0,1],[0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```



```
In [ ]: # aggregating all the reports in one DF
reports_df = {}

models = ["DecisionTree", "RandomForest1", "RandomForest2", "GradientBoosting", "XGBoost", "LightGBM"]
reports = [dt_report, rf_model1, rf_model2, grid_gbd, xgb_model, lgbm_model]

for model, r in zip(models, reports):
    try:
        print(model)
        # t=pd.DataFrame(r)
        # t['model']=model
        temp = pd.DataFrame(r).transpose()
        temp['model']=model
        reports_df[model]=temp
    except:
        continue
    # break
```

```
DecisionTree
RandomForest1
RandomForest2
GradientBoosting
XGBoost
LightGBM
```

```
In [ ]: final_df = pd.concat(reports_df.values()).reset_index()
cols = final_df.columns.tolist()
cols = cols[-1:] + cols[:-1]
final_df = final_df[cols]
```

```
In [ ]: final_df
```

Out[ ]:

	model	index	precision	recall	f1-score	support
0	DecisionTree	0	0.856287	0.934641	0.893750	153.000000
1	DecisionTree	1	0.967742	0.925926	0.946372	324.000000
2	DecisionTree	accuracy	0.928721	0.928721	0.928721	0.928721
3	DecisionTree	macro avg	0.912015	0.930283	0.920061	477.000000
4	DecisionTree	weighted avg	0.931992	0.928721	0.929493	477.000000
5	RandomForest1	0	0.910256	0.928105	0.919094	153.000000
6	RandomForest1	1	0.965732	0.956790	0.961240	324.000000
7	RandomForest1	accuracy	0.947589	0.947589	0.947589	0.947589
8	RandomForest1	macro avg	0.937994	0.942447	0.940167	477.000000
9	RandomForest1	weighted avg	0.947938	0.947589	0.947722	477.000000
10	RandomForest2	0	0.901961	0.901961	0.901961	153.000000
11	RandomForest2	1	0.953704	0.953704	0.953704	324.000000
12	RandomForest2	accuracy	0.937107	0.937107	0.937107	0.937107
13	RandomForest2	macro avg	0.927832	0.927832	0.927832	477.000000
14	RandomForest2	weighted avg	0.937107	0.937107	0.937107	477.000000
15	XGBoost	0	0.875000	0.915033	0.894569	153.000000
16	XGBoost	1	0.958991	0.938272	0.948518	324.000000
17	XGBoost	accuracy	0.930818	0.930818	0.930818	0.930818
18	XGBoost	macro avg	0.916995	0.926652	0.921543	477.000000
19	XGBoost	weighted avg	0.932050	0.930818	0.931213	477.000000
20	LightGBM	0	0.886792	0.921569	0.903846	153.000000
21	LightGBM	1	0.962264	0.944444	0.953271	324.000000
22	LightGBM	accuracy	0.937107	0.937107	0.937107	0.937107
23	LightGBM	macro avg	0.924528	0.933007	0.928559	477.000000
24	LightGBM	weighted avg	0.938056	0.937107	0.937418	477.000000

## Model Performance Metrics

### Decision Tree

- Precision (Class 0): 0.856287
- Recall (Class 0): 0.934641
- F1-score (Class 0): 0.893750
- Precision (Class 1): 0.967742
- Recall (Class 1): 0.925926
- F1-score (Class 1): 0.946372

- Accuracy: 0.928721
- Macro avg Precision: 0.912015
- Macro avg Recall: 0.930283
- Macro avg F1-score: 0.920061
- Weighted avg Precision: 0.931992
- Weighted avg Recall: 0.928721
- Weighted avg F1-score: 0.929493

### **Random Forest 1**

- Precision (Class 0): 0.910256
- Recall (Class 0): 0.928105
- F1-score (Class 0): 0.919094
- Precision (Class 1): 0.965732
- Recall (Class 1): 0.956790
- F1-score (Class 1): 0.961240
- Accuracy: 0.947589
- Macro avg Precision: 0.937994
- Macro avg Recall: 0.942447
- Macro avg F1-score: 0.940167
- Weighted avg Precision: 0.947938
- Weighted avg Recall: 0.947589
- Weighted avg F1-score: 0.947817

### **Random Forest 2**

- Precision (Class 0): 0.901961
- Recall (Class 0): 0.901961
- F1-score (Class 0): 0.901961
- Precision (Class 1): 0.953704
- Recall (Class 1): 0.953704
- F1-score (Class 1): 0.953704
- Accuracy: 0.937107
- Macro avg Precision: 0.927832
- Macro avg Recall: 0.927832
- Macro avg F1-score: 0.927832
- Weighted avg Precision: 0.937107
- Weighted avg Recall: 0.937107
- Weighted avg F1-score: 0.937107

### **XGBoost**

- Precision (Class 0): 0.875000
- Recall (Class 0): 0.915033
- F1-score (Class 0): 0.894584
- Precision (Class 1): 0.958991
- Recall (Class 1): 0.938272
- F1-score (Class 1): 0.948518

- Accuracy: 0.930818
- Macro avg Precision: 0.916995
- Macro avg Recall: 0.926652
- Macro avg F1-score: 0.921451
- Weighted avg Precision: 0.932050
- Weighted avg Recall: 0.930818
- Weighted avg F1-score: 0.931213

### LightGBM

- Precision (Class 0): 0.886792
- Recall (Class 0): 0.915690
- F1-score (Class 0): 0.900846
- Precision (Class 1): 0.962264
- Recall (Class 1): 0.944444
- F1-score (Class 1): 0.953256
- Accuracy: 0.937107
- Macro avg Precision: 0.924528
- Macro avg Recall: 0.930067
- Macro avg F1-score: 0.927051
- Weighted avg Precision: 0.938056
- Weighted avg Recall: 0.937107
- Weighted avg F1-score: 0.937148

### Model Selection

For a ride-sharing platform focused on driver churn, the key metric depends on the business objectives:

1. High Recall (Class 1): If the goal is to catch as many potential churners as possible, recall is crucial. Random Forest 1 has the highest recall for Class 1 (0.956790).
2. Balanced Performance (F1-score): The F1-score provides a balance between precision and recall. Random Forest 1 has the highest F1-score for Class 1 (0.961240).
3. Accuracy: Overall performance across all classes can be reflected by accuracy. Random Forest 1 has the highest accuracy (0.947589).
4. Weighted Metrics: Considering the weighted averages of precision, recall, and F1-score gives an overview of the model's performance across all classes. Random Forest 1 performs best with weighted avg F1-score (0.947817).

### Conclusion

Random Forest 1 stands out as the best model based on these metrics:

- It has the highest recall and F1-score for Class 1, which is important for identifying potential churners.
- It has the highest accuracy and balanced performance across all classes.

## Actionable Insights and Recommendation

- based on the data we saw that ~68% of the drivers have left the company.
- after training the model, best features to represent our data are months worked in the company , total business value generated , increase of quarterly rating , uareterly ratings and monthly average income .
- comapny needs to employ more insentive or the drivers based on their worked monthsm business value they generate.
- also company can look at doing more frrequent rating of drivers instead of quarterly rating to see how drivers are performing and provide insentives/bonus to top performers.
- Company can also employ a feedback system where drivers can express their intent for leaving, that way company can promptly see what they can do from their based for that particular driver to make him.her stay

In [ ]: